



# Bayesian forecasting in economics and finance: A modern review<sup>☆</sup>

Gael M. Martin<sup>a,\*</sup>, David T. Frazier<sup>a</sup>, Worapree Maneesoonthorn<sup>a</sup>,  
Rubén Loaiza-Maya<sup>a</sup>, Florian Huber<sup>b</sup>, Gary Koop<sup>c</sup>, John Maheu<sup>d</sup>,  
Didier Nibbering<sup>a</sup>, Anastasios Panagiotelis<sup>e</sup>

<sup>a</sup> Monash University, Australia

<sup>b</sup> University of Salzburg, Austria

<sup>c</sup> University of Strathclyde, UK

<sup>d</sup> McMaster University, Canada

<sup>e</sup> University of Sydney, Australia

## ARTICLE INFO

### Keywords:

Bayesian prediction  
Macroeconomics  
Finance  
Marketing  
Electricity demand  
Bayesian computational methods  
Loss-based Bayesian prediction

## ABSTRACT

The Bayesian statistical paradigm provides a principled and coherent approach to probabilistic forecasting. Uncertainty about all unknowns that characterize any forecasting problem – model, parameters, latent states – is able to be quantified explicitly and factored into the forecast distribution via the process of integration or averaging. Allied with the elegance of the method, Bayesian forecasting is now underpinned by the burgeoning field of Bayesian computation, which enables Bayesian forecasts to be produced for virtually any problem, no matter how large or complex. The current state of play in Bayesian forecasting in economics and finance is the subject of this review. The aim is to provide the reader with an overview of modern approaches to the field, set in some historical context, with sufficient computational detail given to assist the reader with implementation.

© 2023 International Institute of Forecasters. Published by Elsevier B.V. All rights reserved.

## Contents

1. Introduction.....	812
1.1. Why Bayesian forecasting?.....	812
1.2. The purview of this review.....	813
2. A tutorial on Bayesian forecasting.....	813
2.1. The Bayesian forecasting method.....	813
2.2. An overview of computation.....	814
3. Bayesian forecasting: A chronological tour.....	815
3.1. The late 20th century: The advent of MCMC.....	815
3.1.1. Gibbs sampling.....	815
3.1.2. MH-within-Gibbs sampling.....	816

<sup>☆</sup> The authors would like to thank the Editor, an associate editor and two anonymous referees for very helpful and constructive comments on an earlier draft of the paper. Martin, Frazier and Maneesoonthorn have been supported by Australian Research Council Discovery Grant DP200101414. Frazier has also been supported by Australian Research Council Discovery Early Career Researcher Award DE200101070, and Loaiza-Maya supported by Australian Research Council Discovery Early Career Researcher Award DE230100029. Huber has been supported by Austrian Science Fund Grant ZK-35. Maheu is grateful for financial support from the Social Sciences and Humanities Research Council of Canada.

\* Corresponding author.

E-mail address: [gael.martin@monash.edu](mailto:gael.martin@monash.edu) (G.M. Martin).

3.1.3.	MCMC, data augmentation, and state-space models .....	816
3.2.	The 21st century: Intractable forecasting models .....	817
3.2.1.	What do we mean by ‘intractable’? .....	817
3.2.2.	Exact computational solutions .....	817
3.2.3.	Approximate computational solutions .....	818
3.3.	The 21st Century: Misspecified forecasting models .....	819
3.3.1.	The role of model specification in Bayesian forecasting .....	819
3.3.2.	Focused (or ‘loss-based’) Bayesian prediction .....	819
3.3.3.	Bayesian predictive combinations: Beyond BMA .....	820
3.3.4.	Bayesian predictive decision synthesis .....	820
4.	Selective discipline-specific reviews of Bayesian forecasting .....	821
4.1.	Macroeconomics (Florian Huber and Gary Koop) .....	821
4.2.	Finance (John Maheu, Worapree Maneesoonthorn, and Gael Martin) .....	824
4.3.	Marketing (Rubén Loaiza-Maya and Didier Nibbering) .....	826
4.4.	Electricity pricing and demand (Anastasios Panagiotelis) .....	828
5.	In summary .....	830
	Declaration of competing interest .....	830
	Appendix. Further computational details .....	830
A.1.	Gibbs sampling .....	830
A.2.	MH-within-Gibbs sampling .....	830
A.3.	MH-within-Gibbs sampling in state-space models .....	830
A.4.	PMMH in state-space models .....	830
A.5.	ABC based on summary statistics .....	831
A.6.	BSL based on summary statistics .....	831
A.7.	VB .....	831
A.8.	INLA .....	832
	References .....	833

## 1. Introduction

### 1.1. Why Bayesian forecasting?

The Bayesian statistical paradigm uses the rules and language of probability to quantify uncertainty about all unknown aspects of phenomena that generate observed data. This core characteristic of the paradigm makes it particularly suitable for forecasting, with uncertainty about the unknown values of future observations automatically expressed in terms of a probability distribution. Moreover, Bayesian methods – in principle – allow a user to seamlessly and systematically yield probabilistic forecasts that reflect uncertainty about all unknowns and that, as a consequence, condition primarily on *known* past events or data – a feature that Geweke and Whiteman (2006) refer to as the *principle of relevant conditioning*.

Indeed, the ability of Bayesian forecasters to appropriately incorporate the uncertainty associated with the production of forecasts, while utilizing all available information (both *a priori* and sample information) in a principled manner, led Granger et al. (1986) to conclude that:

“In terms of forecasting accuracy a good Bayesian will beat a non-Bayesian, who will do better than a bad Bayesian.”

Echoing these sentiments, in our opinion, the power of the Bayesian forecasting paradigm is a product of the paradigm’s ability to treat all elements of the statistical problem necessary to produce forecasts – future observations, past observations, parameters, latent variables,

models – as arguments of a joint probability distribution. The express probabilistic formulation of these elements, in turn, allows a Bayesian to invoke the standard rules of probability to produce a distribution for an unknown future value that is conditioned on the known past data, and is *marginal* of other arguments that are inherently unknown.

While this ability to marginalize all unknowns through probability calculus is the hallmark of the Bayesian approach, the benefit of the paradigm, and what ultimately in our opinion defines a ‘good Bayesian’, is the attention to detail necessary to successfully implement Bayesian methods. In Bayesian forecasting, before we ever attempt to produce a forecast, we must first carefully enumerate all possible sources of uncertainty – including, where possible, the set of alternative forecasting models – and construct reasonable prior beliefs for these quantities, which often include (possibly several layers of) latent variables that have a specific and delicate interaction with the observed data, always taking great care to ensure that these prior beliefs do not conflict with the observed data. Then and only then can we ‘turn the Bayesian crank’ to produce the joint posterior distribution over all unknown quantities (including future values), and ultimately integrate out the quantities we are not interested in to obtain the (posterior) predictive distribution for the future values of our random variables of interest. The attention to detail necessary to produce Bayesian forecasts aims to reduce the number of implicit maintained assumptions, and what explicit assumptions are maintained (e.g. the conditioning on a particular model or finite model set) can often be rationalized/tested against the data.

Consistent with the internal coherence of the Bayesian statistical paradigm, the basic manner in which all Bayesian

forecasting problems are *framed* is the same. What differs from case to case, however, is the way in which the problem is *solved* – i.e. the way in which the forecast distribution is accessed. To understand why this is so, it is sufficient to recognize that virtually all Bayesian quantities of interest, including forecast distributions, can be expressed as expectations of some sort. For most models that are used to predict empirically relevant data, these expectations are not available in closed form. Hence, in any practical problem, the implementation of Bayesian forecasting is both model- and data-dependent, and relies on advanced computational tools. Different forecasting problems – defined by different forms and sizes of models and data sets – require, in turn, different approaches to computation. The evolution of the practice of Bayesian forecasting has, as a consequence, gone hand-in-hand with developments on the computational front, with increasingly large and complex models rendered amenable to a Bayesian forecasting approach via access to modern techniques of computation.

## 1.2. The purview of this review

In this review, we give a modern take on the current landscape of Bayesian forecasting. Whilst excellent textbook treatments of Bayesian forecasting are given in Geweke (2005) and West and Harrison (2006), and with Geweke and Whiteman (2006) reviewing specific aspects of Bayesian forecasting in a slew of practical settings, the field has advanced by leaps and bounds in the last 20 years. Therefore, we believe the time is ripe to consider a review of the subject that touches on many of the novel and exciting areas now being explored. The methodological advances we review have general applicability to all discipline areas. Nevertheless, due to our own interests, expertise, and experience – and to keep the scope of the paper manageable – we have chosen to focus primarily on applications in the economic sciences. Whilst the paper is not designed to be a treatise on Bayesian computation, sufficient details are provided to enable the practitioner to understand *why* numerical tools are needed in most forecasting settings, and *how* they are used.

The general structure of the paper is as follows. In Section 2 we provide a short tutorial on Bayesian forecasting. This begins with an outline of the Bayesian forecasting method, followed by an overview of the computational techniques used to implement the method. In Section 3 we then take the reader on a potted chronological tour of Bayesian forecasting, up to the present day. We begin by giving a snapshot of the forecasting problems tackled during the last decade of the 20th century (and the early years of the 21st), and the computational solutions that were adopted then, most notably, Markov chain Monte Carlo (MCMC) algorithms. We then look at the types of intractable forecasting problems that are increasingly encountered in the 21st century, and provide an overview of the new computational solutions that have been proposed to tackle such problems. We also outline very recent developments in which misspecification of the forecasting model is explicitly acknowledged, and conventional

likelihood-based Bayesian forecasting eschewed as a consequence, with problem-specific measures of forecast accuracy (or forecast loss) used, instead, to drive the production of forecast distributions. Section 4 then provides the reader with more detailed reviews of contemporary Bayesian forecasting in the following four broad fields: macroeconomics, finance, marketing, and electricity pricing and demand. Section 5 closes the paper with a brief summary of the current state of play.

Before proceeding further, we make a note about scope and language. To render the scope of the paper manageable, we focus primarily on Bayesian forecasting in time series models – i.e. models for random variables that are indexed by time – and on using such models to say something about the values that these random variables will assume in the future. These future values may be informed only by past observations on the variable, or may also depend on the known values of covariates or regressors. We also follow the convention in the Bayesian literature by using the terms ‘forecast’ and ‘prediction’ (and all of their various grammatical derivations) synonymously and interchangeably in this case, for the sake of linguistic variety. The fundamental principles of Bayesian prediction apply equally to data indexed by something other than time. The term ‘forecast’ is not used in this case, as it is a term reserved for temporal settings. The main exceptions to our focus on time series models and forecasting *per se* occur in Section 4.3, in which models for cross-sectional data are used to predict customer choice in marketing settings, and Section 4.4, in which models for electricity demand that have a spatial dimension are referenced.

## 2. A tutorial on Bayesian forecasting

### 2.1. The Bayesian forecasting method

For the sake of illustration, we assume a scalar random variable  $y_t$ , and define the  $(T \times 1)$  vector of observations on  $y_t$  as  $\mathbf{y}_{1:T} = (y_1, y_2, \dots, y_T)'$ . We assume (for the moment) that  $\mathbf{y}_{1:T}$  has been generated from some parametric model with likelihood  $p(\mathbf{y}_{1:T}|\boldsymbol{\theta})$ , with  $\boldsymbol{\theta} = (\theta_1, \theta_2, \dots, \theta_p)' \in \Theta \subseteq \mathbb{R}^p$  a  $p$ -dimensional vector of unknown parameters, and where we possess prior beliefs on  $\boldsymbol{\theta}$  specified by  $p(\boldsymbol{\theta})$ . Using the same symbol  $\mathbf{y}_{1:T}$  to denote both the vector of observed data and the  $T$ -dimensional vector random variable, we define the joint distribution over  $\mathbf{y}_{1:T}$  and  $\boldsymbol{\theta}$  as  $p(\mathbf{y}_{1:T}, \boldsymbol{\theta})$ . The application of the standard rules of probability to  $p(\mathbf{y}_{1:T}, \boldsymbol{\theta})$  yields *Bayes' theorem* (or *Bayes' rule*):

$$p(\boldsymbol{\theta}|\mathbf{y}_{1:T}) = \frac{p(\mathbf{y}_{1:T}|\boldsymbol{\theta})p(\boldsymbol{\theta})}{p(\mathbf{y}_{1:T})}, \quad (1)$$

where  $p(\mathbf{y}_{1:T}) = \int_{\Theta} p(\mathbf{y}_{1:T}|\boldsymbol{\theta})p(\boldsymbol{\theta})d\boldsymbol{\theta}$ . Bayes' theorem provides a representation for the posterior probability density function (pdf) for  $\boldsymbol{\theta}$ ,  $p(\boldsymbol{\theta}|\mathbf{y}_{1:T})$ , as proportional to the product of the likelihood function and the prior. The term  $p(\mathbf{y}_{1:T})$  defines the marginal likelihood, and the scale factor  $[p(\mathbf{y}_{1:T})]^{-1}$  in (1) ensures that  $p(\boldsymbol{\theta}|\mathbf{y}_{1:T})$  integrates to one.

Now, define  $y_{T+1}$  as the (one-step-ahead) future random variable, where we focus on one-step-ahead forecasting in Sections 2 and 3 merely to simplify the exposition. Assuming  $y_{T+1}$  to be a continuous random variable (again, for illustration), standard probability manipulations lead to the following expression for the *forecast* (or *predictive*) pdf for  $y_{T+1}$ :

$$p(y_{T+1}|\mathbf{y}_{1:T}) = \int_{\Theta} p(y_{T+1}|\theta, \mathbf{y}_{1:T})p(\theta|\mathbf{y}_{1:T})d\theta. \quad (2)$$

When no confusion arises, we also refer to  $p(y_{T+1}|\mathbf{y}_{1:T})$ , albeit loosely, as the forecast (or predictive) *distribution*, or simply as the ‘predictive’.<sup>1</sup> The density  $p(y_{T+1}|\mathbf{y}_{1:T})$  summarizes all uncertainty about  $y_{T+1}$ , conditional on the assumed model – which underpins the structure of both the conditional predictive,  $p(y_{T+1}|\theta, \mathbf{y}_{1:T})$ , and the posterior itself – and the prior beliefs that inform  $p(\theta|\mathbf{y}_{1:T})$ . Point and interval predictions of  $y_{T+1}$ , and indeed any other distributional summary, can be extracted from (2). In the case where the model itself is uncertain, and a finite set of parametric models,  $\mathcal{M}_1, \mathcal{M}_2, \dots, \mathcal{M}_K$ , is assumed to span the model space, a model-averaged predictive (e.g. Raftery, Madigan, & Hoeting, 1997; Section 2),  $p_{MA}(y_{T+1}|\mathbf{y}_{1:T})$ , is produced as

$$p_{MA}(y_{T+1}|\mathbf{y}_{1:T}) = \sum_{k=1}^K p(y_{T+1}|\mathbf{y}_{1:T}, \mathcal{M}_k)p(\mathcal{M}_k|\mathbf{y}_{1:T}), \quad (3)$$

where  $p(y_{T+1}|\mathbf{y}_{1:T}, \mathcal{M}_k)$  denotes the density in (2), but now conditioned explicitly on the  $k$ th model in the set. The  $k$ th posterior model probability,  $p(\mathcal{M}_k|\mathbf{y}_{1:T})$ ,  $k = 1, 2, \dots, K$ , is computed via a further application of Bayes’ theorem in which the (initial) joint distribution of interest is defined over both the model space and the space for the parameters of each of the  $K$  models. Standard manipulations lead to

$$p(\mathcal{M}_k|\mathbf{y}_{1:T}) \propto p(\mathbf{y}_{1:T}|\mathcal{M}_k)p(\mathcal{M}_k), \quad (4)$$

where

$$p(\mathbf{y}_{1:T}|\mathcal{M}_k) = \int_{\Theta_k} p(\mathbf{y}_{1:T}|\theta_k, \mathcal{M}_k)p(\theta_k|\mathcal{M}_k)d\theta_k, \quad (5)$$

for each  $k = 1, 2, \dots, K$ , with  $\theta_k$  denoting the parameter set for the  $k$ th model.

As is clear, *analytical* evaluation of  $p(y_{T+1}|\mathbf{y}_{1:T})$  in (2) requires, at the very least, a closed-form expression for  $p(\theta|\mathbf{y}_{1:T})$ . Typically, however, such an expression is not available, with most posteriors being known only up to a constant of proportionality, as

$$p(\theta|\mathbf{y}_{1:T}) \propto p(\mathbf{y}_{1:T}|\theta)p(\theta). \quad (6)$$

<sup>1</sup> We note that  $p(y_{T+1}|\mathbf{y}_{1:T})$  is sometimes referred to as a ‘posterior’ predictive in the literature, given that it is produced by averaging the conditional predictive,  $p(y_{T+1}|\theta, \mathbf{y}_{1:T})$ , with respect to the posterior density,  $p(\theta|\mathbf{y}_{1:T})$ . We do not adopt this expression, leaving it to the context to make it clear as to whether the term ‘predictive’ is being used to refer to the distribution that is marginal of  $\theta$ ,  $p(y_{T+1}|\mathbf{y}_{1:T})$ , or that which is conditioned on  $\theta$ ,  $p(y_{T+1}|\theta, \mathbf{y}_{1:T})$ . We also streamline the exposition by not using explicit notation for any observed covariates on which the model for  $y_t$  may depend, and on which the predictive for  $y_{T+1}$  would condition, unless this is essential.

The main exceptions to this occur when  $p(\mathbf{y}_{1:T}|\theta)$  is from the exponential family, and either a natural conjugate or a convenient non-informative prior is adopted. Such specifications may be suitable for some simple (and low-dimensional) empirical problems, but are certainly not broadly applicable in practice. Analytical evaluation of  $p_{MA}(y_{T+1}|\mathbf{y}_{1:T})$  in (3) also requires a closed-form expression for each  $p(\mathbf{y}_{1:T}|\mathcal{M}_k)$  (with normalization of  $p(\mathcal{M}_k|\mathbf{y}_{1:T})$  then being straightforward); once again a rare thing beyond the exponential family (and standard prior) setting. Hence, numerical *computation* is needed to implement Bayesian forecasting in virtually all realistic empirical problems.<sup>2</sup>

## 2.2. An overview of computation

The form of (2) makes it clear that the Bayesian predictive pdf,  $p(y_{T+1}|\mathbf{y}_{1:T})$ , is nothing more than the posterior expectation of the predictive that conditions on  $\theta$ . Hence, accessing  $p(y_{T+1}|\mathbf{y}_{1:T})$  amounts to the evaluation of an expectation. This insight is helpful, as it enables us to see many of the computational methods that are used to access  $p(y_{T+1}|\mathbf{y}_{1:T})$  – in cases where it is not available in closed form – simply as different ways of numerically estimating an expectation.

It is convenient to group Bayesian computational methods into three categories: (i) deterministic integration (or quadrature) methods (Davis & Rabinowitz, 1975; Naylor & Smith, 1982), (ii) exact simulation methods, and (iii) approximate methods. Given that the production of  $p(y_{T+1}|\mathbf{y}_{1:T})$  involves integration over  $\theta$ , only in very low-dimensional models is (i) a feasible computational approach on its own, due to the well-known curse of dimensionality that characterizes numerical quadrature. Hence, the computational methods in (ii) and (iii) are those most commonly adopted, and will be our focus here. Note that quadrature may still play a limited role *within* these alternative computational frameworks.

The methods in (ii) use simulation to produce  $M$  draws of  $\theta$ ,  $\theta^{(i)}$ ,  $i = 1, 2, \dots, M$ , from the posterior  $p(\theta|\mathbf{y}_{1:T})$ , which in turn define  $M$  conditional predictives,  $p(y_{T+1}|\theta^{(i)}, \mathbf{y}_{1:T})$ ,  $i = 1, 2, \dots, M$ , the mean of which is used to estimate (2). Alternatively, if it is easier to simulate from  $p(y_{T+1}|\theta^{(i)}, \mathbf{y}_{1:T})$  than to evaluate it at any point in the support of  $y_{T+1}$ ,  $M$  draws of  $y_{T+1}$ ,  $y_{T+1}^{(i)}$ ,  $i = 1, 2, \dots, M$ , are taken, one for each draw  $\theta^{(i)}$ , and kernel density estimation methods are used to produce an estimate of  $p(y_{T+1}|\mathbf{y}_{1:T})$ . Different simulation methods are distinguished by the way in which the posterior draws are produced. Methods in (ii) include Monte Carlo sampling (Metropolis & Ulam, 1949), importance sampling (IS) (Hammersley & Handscomb, 1964; Kloek & van Dijk, 1978; Geweke, 1989), and MCMC sampling – including Gibbs sampling (Geman & Geman, 1984; Gelfand & Smith, 1990) and Metropolis–Hastings (MH) algorithms

<sup>2</sup> Numerous textbook illustrations of the material in this section can be found. In addition to Geweke (2005) and West and Harrison (2006) as cited above, some examples are Zellner (1971), Koop (2003), and Robert (2007). We also refer the reader to Steel (2020) for a recent review of Bayesian model averaging in economics.



(Metropolis, Rosenbluth, Rosenbluth, Teller, & Teller, 1953; Hastings, 1970) – with MCMC being by far the most common simulation method used to compute forecast distributions in practice. The term ‘exact’ arises from the fact that, under appropriate conditions (including convergence of the Markov chain to  $p(\theta|\mathbf{y}_{1:T})$  in the case of the MCMC algorithms), such methods all produce a  $\sqrt{M}$ -consistent estimate of the ordinate  $p(y_{T+1}|\mathbf{y}_{1:T})$  at any point in the support of the random variable  $y_{T+1}$ ; this estimate can thus be rendered arbitrarily accurate for large enough  $M$ .

We refer the reader to Chib and Greenberg (1996) and Geyer (2011) for reviews of MCMC sampling; to Casella and George (1992) and Chib and Greenberg (1995) for descriptions of the Gibbs and MH algorithms, respectively, that are useful for practitioners; and to Andrieu, Doucet, and Robert (2004), Robert and Casella (2011), and Martin, Frazier, and Robert (2023b) for historical accounts of MCMC sampling. Geweke and Whiteman (2006) also serves as an excellent reference on the use of these computational methods in a forecasting context. Given the critical role played by MCMC methods in the production of Bayesian forecasts, the basic principles of the algorithms are also outlined below in Section 3.1, with more recent developments of both IS and MCMC – most notably sequential Monte Carlo (SMC) (Gordon, Salmond, & Smith, 1993; Chopin & Papaspiliopoulos, 2020) and pseudo-marginal MCMC (Beaumont, 2003; Andrieu & Roberts, 2009; Andrieu, Doucet, & Holenstein, 2011) – discussed briefly in Section 3.2.

The methods in (iii) replace  $p(\theta|\mathbf{y}_{1:T})$  in the integrand of (2) with an *approximation* of some sort, and evaluate the resultant integral. In so doing, such methods do not aim to estimate  $p(y_{T+1}|\mathbf{y}_{1:T})$  itself, but rather some representation of it, defined as the expectation of  $p(y_{T+1}|\theta, \mathbf{y}_{1:T})$  with respect to the relevant posterior approximation. The methods in (iii) have been based on the principles of approximate Bayesian computation (ABC) (Marin, Pudlo, Robert, & Ryder, 2011; Sisson & Fan, 2011; Sisson, Fan, & Beaumont, 2019), Bayesian synthetic likelihood (BSL) (Price, Drovandi, Lee, & Nott, 2018), variational Bayes (VB) (Blei, Kucukelbir, & McAuliffe, 2017), and integrated nested Laplace approximation (INLA) (Rue, Martino, & Chopin, 2009), and they produce what are termed ‘approximate’ forecast, or predictive distributions. It suffices to say that the principle adopted for estimating the approximate predictive so-defined is typically one and the same: Draws of  $\theta$  from the approximate posterior (however produced) are used to produce either a sample mean of conditional predictives or  $M$  draws of  $y_{T+1}$  from  $p(y_{T+1}|\theta, \mathbf{y}_{1:T})$ , with kernel density estimation then applied.

The production of (3) requires the computation of each model-specific predictive, plus the computation of each (5). The first set of  $K$  computations would proceed via the sorts of steps outlined above. Computation of the  $K$  marginal likelihoods could also be performed via one of the three broad methods listed above, in particular (ii) or (iii). However, the fact that each (5) is a prior, rather than a posterior expectation does have implications for the precise manner in which the computation is implemented. See Ardia, Baştürk, Hoogerheide, & van Dijk, 2012 and Llorente, Martino, Delgado, & Lopez-Santiago, 2021 for details.

### 3. Bayesian forecasting: A chronological tour

#### 3.1. The late 20th century: The advent of MCMC

As is clear from the brief synopsis above, *simulation* is key to computing forecast distributions when they are not available in closed form. While the use of simulation to compute statistical quantities of interest was known by the 1970s (Metropolis & Ulam, 1949; Metropolis et al., 1953; Hammersley & Handscomb, 1964; Hastings, 1970), the technology required to perform simulation in a convenient and timely fashion was not yet available, and simulation-based computation thus remained largely out of reach. To quote Geweke and Whiteman (2006):

“In the beginning, there was diffuseness, conjugacy and analytical work!”

In the latter part of the 20th century, things changed. The increased speed and availability of desktop machines (Ceruzzi, 2003), allied with critical advances in simulation methodology, led to a proliferation of methods for accessing  $p(y_{T+1}|\mathbf{y}_{1:T})$  via the simulation of draws from  $p(\theta|\mathbf{y}_{1:T})$ . To this end, we give a brief outline of the preeminent posterior simulation algorithms of the 1990s (and the early 2000s): Gibbs sampling (Section 3.1.1), MH-within-Gibbs sampling (Section 3.1.2), and (MH-within-) Gibbs sampling allied with data augmentation (Section 3.1.3) – touching on the types of forecasting models that were able to be treated via such methods, most notably the ubiquitous state-space models that underpin much modern Bayesian forecasting. To keep the exposition concise, we place all algorithmic details in the [Appendix](#) and reference specific algorithms from the [Appendix](#) at suitable points in the text.

##### 3.1.1. Gibbs sampling

As a general rule, if  $p(\theta|\mathbf{y}_{1:T})$  does not have a closed-form representation, it is also not amenable to Monte Carlo sampling, as the latter requires that  $p(\theta|\mathbf{y}_{1:T})$  can be decomposed into recognizable densities, from which computer simulation is feasible. IS (Kloek & van Dijk, 1978; Geweke, 1989), via use of an ‘importance’ or ‘proposal’ density,  $q(\theta|\mathbf{y}_{1:T})$ , that matches  $p(\theta|\mathbf{y}_{1:T})$  well and which *can* be drawn from, is a possible solution in some cases. However, the algorithm can fail to produce representative draws from  $p(\theta|\mathbf{y}_{1:T})$  when the dimension of  $\theta$  is large, due to the difficulty of finding a  $q(\theta|\mathbf{y}_{1:T})$  that is a good match to  $p(\theta|\mathbf{y}_{1:T})$  in high dimensions.

In contrast, under certain conditions, a Gibbs sampler is able to produce a (dependent) set of draws from the joint posterior via iterative sampling from lower dimensional, and often standard, conditional posteriors. In other words, a Gibbs sampler takes advantage of the fact that, while joint and marginal posterior distributions are usually complex in form and unable to be simulated from directly, conditional posteriors are often standard and amenable to simulation. Given the satisfaction of the required convergence conditions (Geyer, 2011), draws  $\theta^{(i)}$ ,  $i = 1, 2, \dots, M$ , produced via iterative sampling from the full conditionals, converge in distribution to  $p(\theta|\mathbf{y}_{1:T})$  as

$M \rightarrow \infty$ , and can be used to produce a  $\sqrt{M}$ -consistent estimate of the ordinates of  $p(y_{T+1}|\mathbf{y}_{1:T})$  across the support of  $y_{T+1}$  in the manner described in Section 2.2. Decisions about how to partition (or ‘block’)  $\theta$  need to be made (Liu, Wong, & Kong, 1994; Roberts & Sahu, 1997), with a view to increasing the efficiency of the chain. In effect, this amounts to ensuring an accurate estimate of  $p(y_{T+1}|\mathbf{y}_{1:T})$  for a given number of draws,  $M$ . See Algorithm 1 in Appendix A.1.

Chib (1993) and McCulloch and Tsay (1994) are the earliest examples of using Gibbs algorithms for Bayesian estimation and prediction in time series settings. Both papers focus on models that include an autoregressive (AR) component as part of their structure, and exploit the fact that despite  $p(\theta|\mathbf{y}_{1:T})$  and  $p(y_{T+1}|\mathbf{y}_{1:T})$  precluding analytical treatment in most of the examples considered, the conditional posteriors always have closed forms. As one would anticipate, however, a ‘pure’ Gibbs algorithm based on a full set of standard conditionals is not always possible, with the more typical situation being one in which one or more of the conditionals – associated with any given partitioning of the parameter space – are not available in closed form. The following section describes how to adapt a Gibbs algorithm in cases where certain conditional components are not known in closed form, and, in so doing, illustrates a powerful simulation-based algorithm for accessing  $p(y_{T+1}|\mathbf{y}_{1:T})$  in more complex settings.

### 3.1.2. MH-within-Gibbs sampling

The Gibbs sampler is only one example of an MCMC algorithm. The first such example – the Metropolis algorithm – appeared in a paper that has assumed an important status in the history of statistics: Metropolis et al. (1953).<sup>3</sup> The Metropolis algorithm was subsequently generalized by Hastings (1970), and it is this MH version of the method that is typically referenced. For the purpose of this review, the key role of the MH algorithm is to enable sampling from non-standard conditionals within a Gibbs algorithm, in particular when the dimension of the conditionals precludes, say, the exclusive use of inverse cumulative distribution function (ICDF) sampling.<sup>4</sup>

Under regularity, a Markov chain that converges to  $p(\theta|\mathbf{y}_{1:T})$  can be produced by embedding an MH algorithm (or MH algorithms) within an outer Gibbs loop. In short, an MH-within-Gibbs algorithm proceeds by drawing from any non-standard conditional indirectly, via a ‘candidate’ (or ‘proposal’) distribution that is deemed to be a good

match to the inaccessible conditional, and accepting the draw with a given probability. Critically, the formula that defines the acceptance probability involves evaluating the non-standard conditional only up to its integrating constant; hence, the conditional need not be known in its entirety.<sup>5</sup> Again, under appropriate regularity, the draws  $\theta^{(i)}$ ,  $i = 1, 2, \dots, M$ , from the MH-within-Gibbs algorithm converge in distribution to  $p(\theta|\mathbf{y}_{1:T})$  as  $M \rightarrow \infty$ , and can be used to produce a  $\sqrt{M}$ -consistent estimate of the ordinates of  $p(y_{T+1}|\mathbf{y}_{1:T})$ . See Algorithm 2 in Appendix A.2.

As will become evident in the subsequent empirical review sections, MH-within-Gibbs algorithms remain the dominant form of method used to sample from posteriors – and to estimate predictive distributions – for time series models for which a convenient partitioning of the parameter space is available, and for which the conditional posteriors are known up to their integrating constants. Hence, we reserve further elaboration on the use of such algorithms in practice until the appropriate points in Section 4.

### 3.1.3. MCMC, data augmentation, and state-space models

For many empirical problems in economics and related fields, a suitable model can be partitioned into two sets: static unknowns  $\theta$ , which are fixed throughout time; and latent data  $\mathbf{z}_{1:T} = (z_1, z_2, \dots, z_T)$ , which vary over time. The latent states may be intrinsic to the model – as in a state-space model – or they may be auxiliary variables introduced purely for the purpose of facilitating posterior sampling. The application of a Gibbs-based MCMC scheme to the joint, or ‘augmented’, set of unknowns  $(\theta, \mathbf{z}_{1:T})$  is often referred to as ‘data augmentation’, in the spirit of Tanner and Wong (1987), and such schemes have enabled the Bayesian analysis of large classes of time series models that would have otherwise been inaccessible.

We illustrate here the basic principles of the approach using a state-space model governed by a measurement density for the observed scalar random variable,  $y_t$ , and a Markov transition density for a scalar state variable,  $z_t$ :

$$p(y_t|z_t, \theta) \quad (7)$$

$$p(z_t|z_{t-1}, \theta). \quad (8)$$

Using the generic notation in (7) and (8), the augmented posterior is

$$p(\theta, \mathbf{z}_{1:T}|\mathbf{y}_{1:T}) \propto p(\mathbf{y}_{1:T}|\mathbf{z}_{1:T}, \theta)p(\mathbf{z}_{1:T}|\theta)p(\theta). \quad (9)$$

In certain cases, the model structure is such that a pure Gibbs scheme can be used to produce draws from  $p(\theta, \mathbf{z}_{1:T}|\mathbf{y}_{1:T})$  and, thus, from  $p(\theta|\mathbf{y}_{1:T})$  – an insight obtained independently by Carter and Kohn (1994) and Frühwirth-Schnatter (1994) for the case of the linear Gaussian state-space model, for example. However, the implementation of such a scheme will by definition

<sup>3</sup> For example, Dongarra and Sullivan (2000) rank the Metropolis algorithm proposed in Metropolis et al. as one of the 10 algorithms “with the greatest influence on the development and practice of science and engineering in the 20th century”.

<sup>4</sup> Any non-standard probability distribution can, in principle, be drawn from using ICDF sampling. The term ‘Griddy-Gibbs sampling’ was first used by Ritter and Tanner (1992) to refer to the use of ICDF sampling to draw from non-standard conditionals in a Gibbs scheme. Given that the method amounts to the use of numerical quadrature, it suffers from the curse of dimensionality, and is thus infeasible for drawing from anything other than very low-dimensional conditionals. See Bauwens and Lubrano (1998) for the application of the Griddy-Gibbs sampler to a generalized autoregressive conditionally heteroskedastic (GARCH) model for financial returns.

<sup>5</sup> Moreover, and in contrast to IS, the requirement to find a well-matched proposal distribution is facilitated by the dimension reduction invoked by breaking down the high-dimensional joint posterior into lower-dimensional conditionals before any proposal distribution needs to be specified.

require both  $p(\theta|\mathbf{z}_{1:T}, \mathbf{y}_{1:T})$  and  $p(\mathbf{z}_{1:T}|\theta, \mathbf{y}_{1:T})$  to have recognizable forms. In more general cases, in which either the measurement or state equation has non-linear and/or non-Gaussian features, the resulting conditionals will not necessarily have a known closed form, which necessitates the addition of MH steps within the outer Gibbs loop. Such a treatment was the method of attack for large classes of models in the 1990s and early 2000s. Relevant contributions here, which include specific treatments of the ubiquitous stochastic volatility (SV) model, are Polson, Carlin, and Stoffer (1992), Jacquier, Polson, and Rossi (1994), Shephard and Pitt (1997), Kim, Shephard, and Chib (1998), Chib, Nadari, and Shephard (2002), Stroud, Müller, and Polson (2003), Chib, Nadari, and Shephard (2006), Strickland, Forbes, and Martin (2006), Omori, Chib, Shephard, and Nakajima (2007), and Strickland, Martin, and Forbes (2008). The reviews of Fearnhead (2011) and Giordani, Pitt, and Kohn (2011) provide more detailed accounts and extensive referencing of this earlier literature.<sup>6</sup> (See also Appendix A.3.)

To conclude, and once again using the generic notation in (7) and (8), once draws have been produced from  $p(\theta, \mathbf{z}_{1:T}|\mathbf{y}_{1:T})$ , the predictive pdf,

$$p(y_{T+1}|\mathbf{y}_{1:T}) = \int_{\mathbf{z}_{T+1}} \int_{\mathbf{z}_{1:T}} \int_{\Theta} p(y_{T+1}|\mathbf{z}_{T+1}, \theta, \mathbf{y}_{1:T}) \times p(\mathbf{z}_{T+1}|\mathbf{z}_T, \theta) p(\theta, \mathbf{z}_{1:T}|\mathbf{y}_{1:T}) d\theta d\mathbf{z}_{1:T} d\mathbf{z}_{T+1}, \quad (10)$$

can be estimated in the usual way, using subsequent draws from  $p(\mathbf{z}_{T+1}|\mathbf{z}_T, \theta)$  and  $p(y_{T+1}|\mathbf{z}_{T+1}, \theta, \mathbf{y}_{1:T})$ , or by averaging the conditional predictives over all draws of  $\mathbf{z}_{T+1}$  and  $\theta$ .

### 3.2. The 21st century: Intractable forecasting models

#### 3.2.1. What do we mean by ‘intractable’?

The MCMC methods that first evolved during the late 20th century continue to serve as the bread and butter of Bayesian forecasting, as will be made evident in Section 4. Nevertheless, more ambitious forecasting problems are now being tackled, and this has tested the mettle of some of the early algorithms. As a consequence, Bayesian forecasters have begun to exploit more modern computational techniques, and it is those techniques that we touch on briefly in this section.

It is convenient to characterize these newer computational developments as different types of solutions to so-called intractable forecasting problems, by which we mean: (a) forecasts based on models with data generating processes (DGPs) that cannot be readily expressed as a pdf or probability mass function (pmf); (b) forecasts based on high-dimensional models, with a very large number of unknowns; and (c) forecasts produced using extremely large data sets. Problems that feature (a) are referred to as ‘doubly intractable problems’, as not only is  $p(\theta|\mathbf{y}_{1:T})$  not

available in its entirety (as is typical), but the DGP itself is also not able to be expressed analytically.

With reference to (a), the MCMC methods referenced so far entail the evaluation of the DGP as a  $pd/(m)f$ , either in the calculation of the acceptance probability in any MH sub-step, or in the specification of full conditionals in any ‘pure’ Gibbs step. Hence, they are infeasible when DGPs do not admit such a representation. Many such DGPs exist (see, for example, Martin, Frazier, & Robert, 2023a; for a list of examples), but particularly pertinent ones to mention here are continuous time models in finance with unknown transition densities (Gallant & Tauchen, 1996),  $\alpha$ -stable models for financial returns (and/or their volatility) (Peters, Sisson, & Fan, 2012; Martin, McCabe, Frazier, Maneesoonthorn, & Robert, 2019), and stochastic dynamic equilibrium models in economics (Calvet & Czellar, 2015). With regard to (b), whilst, in principle (and under appropriate regularity), a convergent MCMC chain can be constructed for any model, the exploration of a very high-dimensional parameter space via an MCMC algorithm can be prohibitively slow (Tavaré, Balding, Griffith, & Donnelly, 1997; Rue et al., 2009; Braun & McAuliffe, 2010; Lintusaari, Gutmann, Dutta, Kaski, & Corander, 2017; Betancourt, 2018; Johndrow, Smith, Pillai, & Dunson, 2019). Hence, in models with a very large number of unknowns – including those with multiple sets of latent variables – the production of an accurate MCMC-based estimate of  $p(y_{T+1}|\mathbf{y}_{1:T})$  in a practical amount of time may not be possible. Finally, regarding (c), MCMC schemes require pointwise (i.e. for each  $y_t$ ) evaluations of  $p(\mathbf{y}_{1:T}|\theta)$  at each draw of  $\theta$ , thereby inducing an  $O(T)$  computational burden at each iteration in an MCMC chain.<sup>7</sup> Such schemes can thus struggle when confronted with ‘big data’ (Bardenet, Doucet, & Holmes, 2017). In this context, the term ‘big data’ refers to situations where, due to the length and/or size of the data set, the repeated evaluation of the likelihood function that is required to produce draws from the corresponding MCMC chain is too time consuming for the algorithm to run in a reasonable amount of time.

The methods in the following sections have been designed to solve one or more of these instances of intractability. The techniques in Section 3.2.2 do so whilst preserving the ‘exact’ nature of the estimate of  $p(y_{T+1}|\mathbf{y}_{1:T})$ , whilst those in Section 3.2.3 aim to produce an approximation of  $p(y_{T+1}|\mathbf{y}_{1:T})$  only.

#### 3.2.2. Exact computational solutions

The first two decades of the 21st century have witnessed a wealth of advances in both MCMC and IS-based algorithms. The goal of the newer MCMC algorithms – at their heart – is to explore the high mass region of the joint posterior more *efficiently*, in particular when the dimension of the space of unknowns is large. This, in turn, enables a more accurate estimate of  $p(y_{T+1}|\mathbf{y}_{1:T})$  to be produced for a given computational budget. This goal has been achieved via a variety of means, which (in the spirit of Robert, Elvira, Tawn, & Wu, 2018 and Martin et al.,

<sup>6</sup> We also note here the work of Chib and Greenberg (1994), in which the state-space representation of an autoregressive moving average (ARMA) model (Harvey, 1981) was exploited, and the principle of data augmentation invoked, in order to enable an MH-within-Gibbs scheme to be applied.

<sup>7</sup> We recall that a sequence  $X_T$  is  $O(T)$  if  $|X_T/T|$  is bounded as  $T \rightarrow +\infty$ .



2023b) can be summarized as follows: the use of more geometric information about the target posterior, most notably the use of Hamiltonian updates (Neal, 2011b; Hoffman & Gelman, 2014); the use of better MH candidate, or proposal distributions, including those that adapt to previous draws (Nott & Kohn, 2005; Roberts & Rosenthal, 2009); various types of combinations of multiple chains (Jacob, Robert, & Smith, 2011; Neal, 2011a; Neiswanger, Wang, & Xing, 2013; Glynn & Rhee, 2014; Huber, 2016; Jacob, O’Leary, & Atchadé, 2020); or the use of *ex post* variance reduction methods (Craiu & Meng, 2005; Douc & Robert, 2011; Owen, 2017; Baker, Fearnhead, Fox, & Nemeth, 2019). We refer the reader to Green, Latuszynski, Pereyra, and Robert (2015), Robert et al. (2018), and Dunson and Johndrow (2020) for detailed reviews of modern developments in MCMC, and to Jahan, Ullah, and Mengersen (2020) for an overview of the way in which certain of the newer methods manage the problem of scale – in terms of either the unknowns or the data, or both.

Whilst not designed expressly to deal with problems of scale, sequential Monte Carlo (SMC) methods – which exploit the principles of IS – have developed in parallel to the expansion of the MCMC stable. Devised initially for the sequential analysis of state-space models, via methods of particle filtering (Gordon et al., 1993), SMC methods have evolved into a larger suite of methods used to perform both sequential and non-sequential tasks (Naesseth, Lindsten, Schön, et al., 2019; Chopin & Papaspiliopoulos, 2020). For the purpose of this review, the most pertinent development is the melding of particle filtering with MCMC in state-space settings to produce a particle marginal MH (PMMH) algorithm (Andrieu et al., 2011; Flury & Shephard, 2011; Pitt, dos Santos Silva, Giordani, & Kohn, 2012; Doucet, Pitt, Deligiannidis, & Kohn, 2015; Deligiannidis, Doucet, & Pitt, 2018). Such algorithms tackle intractability type (a) in the dichotomy of the previous section, by replacing an unavailable likelihood function by an unbiased estimate – produced via the particle filter – in an MH algorithm which, under regularity, retains the posterior  $p(\theta|\mathbf{y}_{1:T})$  as its invariant distribution. Given the increasingly important role played by PMMH, a brief algorithmic description of it is included in Algorithm 3 in Appendix A.4.<sup>8</sup>

### 3.2.3. Approximate computational solutions

In situations in which the dimension or structure of the forecasting model, or the size of the data set, still precludes the use of either an MCMC or a PMMH approach, an approximate method may be the only computational option. The cost of adopting such a solution

is that these methods no longer directly target the exact predictive,  $p(y_{T+1}|\mathbf{y}_{1:T})$ ; instead, an approximation of  $p(y_{T+1}|\mathbf{y}_{1:T})$  becomes the goal.

The spirit of these methods is to approximate  $p(y_{T+1}|\mathbf{y}_{1:T})$  via some feasible approximation to the posterior  $p(\theta|\mathbf{y}_{1:T})$ . Denoting the posterior approximation generically by  $g(\theta|\mathbf{y}_{1:T})$ , the resultant approximate predictive can be expressed as

$$g(y_{T+1}|\mathbf{y}_{1:T}) = \int_{\Theta} p(y_{T+1}|\theta, \mathbf{y}_{1:T}) g(\theta|\mathbf{y}_{1:T}) d\theta, \quad (11)$$

in the case where there are only static unknowns. When the model features both static parameters and time-varying latent parameters, and exploiting the Markov property of the state process in (8), the approximate predictive can be represented as

$$g(y_{T+1}|\mathbf{y}_{1:T}) = \int_{z_{T+1}} \int_{z_T} \int_{\Theta} p(y_{T+1}|z_{T+1}, \theta, \mathbf{y}_{1:T}) \times p(z_{T+1}|z_T, \theta) p(z_T|\theta, \mathbf{y}_{1:T}) g(\theta|\mathbf{y}_{1:T}) d\theta dz_T dz_{T+1}. \quad (12)$$

Given draws of  $\theta$  from  $g(\theta|\mathbf{y}_{1:T})$ , and given an appropriate forward-filtering algorithm to draw from  $p(z_T|\theta, \mathbf{y}_{1:T})$  when needed, a simulation-based estimate of  $g(y_{T+1}|\mathbf{y}_{1:T})$  can be produced in the usual way, either as a sample mean of the conditional predictives defined by the draws of  $\theta$  (and  $z_{T+1}$ ), or by applying kernel density techniques to the draws of  $y_{T+1}$  from the conditional predictive.

With reference to the taxonomy of intractable problems delineated in Section 3.2, the different methods of producing  $g(\theta|\mathbf{y}_{1:T})$  (and, hence,  $g(y_{T+1}|\mathbf{y}_{1:T})$ ) can be categorized according to whether they are being used to obviate (a) or to tackle a problem of scale: (b) and/or (c). Both ABC and BSL avoid the need to evaluate the DGP and, hence, are feasible methods in the doubly intractable settings of category (a). In brief, both methods require only *simulation*, and not *evaluation*, of the DGP. The approximation of  $p(\theta|\mathbf{y}_{1:T})$  arises, primarily, from the fact that both methods (in different ways) degrade the information in the full data set,  $\mathbf{y}_{1:T}$ , to the information contained in a set of summary statistics,  $\eta(\mathbf{y}_{1:T})$ . As such, the target becomes the so-called partial posterior for  $\theta$ , which conditions on  $\eta(\mathbf{y}_{1:T})$ , rather than  $\mathbf{y}_{1:T}$ . The quality of the approximation is thus dependent on the informativeness of the summaries, as well as on other forms of approximation invoked in the implementation of the methods. Vanilla versions of ABC and BSL are provided in Algorithm 4 (Appendix A.5) and Algorithm 5 (Appendix A.6) respectively.

In contrast to ABC and BSL, VB and INLA still target the exact posterior  $p(\theta|\mathbf{y}_{1:T})$ , but provide approximations that can be computationally convenient when the scale of the empirical problem is large in some sense (so problem (b) and/or problem (c)), often as a consequence of the specification of a high number of latent, or ‘local’, parameters in the model, in addition to the (usually) smaller set of ‘global’ parameters ( $\theta$  in our notation). Adopting the technique of the calculus of variations, VB produces an approximation of  $p(\theta|\mathbf{y}_{1:T})$  that is closest to  $p(\theta|\mathbf{y}_{1:T})$  within a chosen variational family, whilst INLA applies a series of nested Laplace approximations

<sup>8</sup> PMMH is actually a special case of the general pseudo-marginal MH technique (also sometimes denoted by the abbreviation ‘PMMH’), in which a pseudo likelihood, produced in some manner or another as an unbiased estimator of the true likelihood, is used within an MH algorithm. See, for example, the subsampling methods based on pseudo-marginal MCMC (Bardenet et al., 2017; Quiroz, Tran, Villani, & Kohn, 2018; Quiroz, Kohn, Villani, & Tran, 2019) used expressly to improve the performance of MCMC in the case of a large-dimensional  $\mathbf{y}_{1:T}$  (i.e. intractability type (c)).



(Laplace, 1774; Tierney & Kadane, 1986; Tierney, Kass, & Kadane, 1989) to a high-dimensional latent Gaussian model to produce an approximation of  $p(\theta|\mathbf{y}_{1:T})$ . Both VB and INLA exploit state-of-the-art optimization techniques, for the purpose of minimizing the distance between  $p(\theta|\mathbf{y}_{1:T})$  and the variational approximation in the case of VB, and for the purpose of producing the mode of the high-dimensional vector of latent states in the case of INLA. The basic principles of VB and INLA are provided in [Appendices A.7](#) and [A.8](#), respectively.

We refer the interested reader to [Martin et al. \(2023a\)](#) for an extensive review of all of these approximate Bayesian methods, as well as more complete coverage of the existing literature, including references to in-depth reviews of specific methods. [Martin et al.](#) also include a discussion of hybrid methods that mix and match features of more than one computational technique, with the aim of tackling multiple instances of intractability simultaneously.

Regardless of which approximation method is used, the hope is that the resulting approximate predictive  $g(y_{T+1}|\mathbf{y}_{1:T})$  performs well relative to the inaccessible exact predictive, and that issue is addressed in certain work cited in the empirical reviews in [Section 4](#).

### 3.3. The 21st Century: Misspecified forecasting models

#### 3.3.1. The role of model specification in Bayesian forecasting

Inherent in the conventional Bayesian approach to forecasting is the assumption that the process that has generated the observed data tallies with the particular model that underpins the likelihood function. Bayesian model averaging (BMA) – and the resultant predictive in [\(3\)](#) – has evolved as a principled way of catering to uncertainty about the predictive model, and BMA remains a very important technique in the Bayesian toolbox. Nevertheless, underpinning BMA is still the assumption that the true process is spanned by the set of models over which one averages – i.e. that the so-called  $\mathcal{M}$ -closed view of the world ([Bernardo & Smith, 1994](#)) prevails.

In response to these perceived limitations of the conventional approach, attention has recently been given to producing predictions that are ‘fit for purpose’, by focusing the Bayesian machinery on the *specific* goals of the predictive analysis at hand. In the following sections we briefly summarize three such approaches, all of which move beyond the conventional likelihood-based Bayesian update, and the  $\mathcal{M}$ -closed paradigm: seeking to produce accurate predictions without recourse to the assumption of correct model specification.

#### 3.3.2. Focused (or ‘loss-based’) Bayesian prediction

[Loaiza-Maya, Martin, and Frazier \(2021\)](#) propose an approach to Bayesian prediction expressly designed for the context of misspecification. In brief, rather than a correct predictive model being assumed, a prior is placed over a class of *plausible* predictive models. The prior is then updated to a posterior via a sample criterion function that is constructed using a scoring rule ([Gneiting & Raftery, 2007](#)) that rewards the type of predictive accuracy (e.g. accurate prediction of extreme values) that is important for the particular empirical problem being

tackled. With a criterion function that explicitly captures predictive accuracy replacing the likelihood function in the Bayesian update, the explicit need for correct model specification is avoided.

Following [Gneiting and Raftery \(2007\)](#), and using generic notation, for  $\mathcal{P}$  a convex class of predictive distributions on  $(\Omega, \mathcal{F})$ , the predictive accuracy of  $P \in \mathcal{P}$  can be assessed using a scoring rule  $S : \mathcal{P} \times \Omega \rightarrow \mathbb{R}$ . If the value  $y$  eventuates, then the positively oriented score of the predictive  $P$  is  $S(P, y)$ . The expected score under the true unknown predictive  $P_0$  is defined as

$$\mathbb{S}(\cdot, P_0) := \int_{y \in \Omega} S(\cdot, y) dP_0(y). \quad (13)$$

A scoring rule is said to be proper relative to  $\mathcal{P}$  if, for all  $P, G \in \mathcal{P}$ ,  $S(G, G) \geq S(P, G)$ , and is strictly proper relative to  $\mathcal{P}$  if  $S(G, G) = S(P, G) \iff P = G$ . Scoring rules are important mechanisms, as they elicit truth telling within the forecasting exercise: if the true predictive  $P_0$  were known, then in terms of forecasting accuracy as measured by the scoring rule  $S(\cdot, \cdot)$ , it would be optimal to use  $P_0$ .

Different scoring rules reward different forms of predictive accuracy (see [Gneiting & Raftery, 2007](#), [Opschoor, Van Dijk, & van der Wel, 2017](#), and [Martin, Loaiza-Maya, Maneesoonthorn, Frazier, & Ramírez-Hassan, 2022](#) for expositions); hence the motivation to drive the update by the score that ‘matters’. Since  $P_0$  and the expected score  $\mathbb{S}(\cdot, P_0)$  are unattainable in practice, an estimate based on  $\mathbf{y}_{1:T}$  is used to define the sample criterion,  $S_T(\theta) := \sum_{t=0}^{T-1} S(p(y_{t+1}|\theta, \mathbf{y}_{1:t}), y_{t+1})$ , where  $p(y_{t+1}|\theta, \mathbf{y}_{1:t})$  is the pdf associated with a given  $P$ . Adopting the exponential updating rule proposed by [Bissiri, Holmes, and Walker \(2016\)](#) (see also [Giummolè, Mameli, Ruli, & Ventura, 2017](#), [Holmes & Walker, 2017](#), [Guedj, 2019](#), [Lyddon, Holmes, & Walker, 2019](#), and [Syring & Martin, 2019](#)), [Loaiza-Maya et al. \(2021\)](#) define the *generalized* (or *Gibbs*) posterior:

$$\pi_w(\theta|\mathbf{y}_{1:T}) = \frac{\exp[wS_T(\theta)]\pi(\theta)}{\int_{\Theta} \exp[wS_T(\theta)]\pi(\theta)d\theta}, \quad (14)$$

for some learning rate  $w \geq 0$ , calibrated in a preliminary step. This posterior explicitly places high weight on – or ‘focuses’ on – values of  $\theta$  that yield high predictive accuracy in the scoring rule  $S(\cdot, \cdot)$ . As such, the process of building a Bayesian predictive as

$$p_w(y_{T+1}|\mathbf{y}_{1:T}) = \int_{\Theta} p(y_{T+1}|\theta, \mathbf{y}_{1:T})\pi_w(\theta|\mathbf{y}_{1:T})d\theta, \quad (15)$$

is termed ‘focused Bayesian prediction’ (FBP) by the authors. By construction, when the predictive model,  $p(y_{T+1}|\theta, \mathbf{y}_{1:T})$ , is misspecified, [\(15\)](#) will – out-of-sample – often outperform, in the chosen rule  $S(\cdot, \cdot)$ , the likelihood (or log-score)-based predictive in [\(2\)](#), and this is demonstrated in [Loaiza-Maya et al.](#) both theoretically and in extensive numerical illustrations.

Since a positively oriented score can, equivalently, be viewed as the negative of a measure of predictive loss, FBP can also be referred to as ‘loss-based’ prediction. Such terminology is indeed adopted in [Frazier, Loaiza-Maya, Martin, and Koo \(2022\)](#), in which the principles delineated here are extended to high-dimensional models,

and approximations to both  $\pi_w(\theta|\mathbf{y}_{1:T})$  and  $p_w(y_{T+1}|\mathbf{y}_{1:T})$  based on VB are proposed and validated. We note that the term ‘loss’ as it is used in [Loaiza-Maya et al. \(2021\)](#) and [Frazier, Loaiza-Maya, Martin, and Koo \(2022\)](#) refers specifically to predictive loss as quantified by a proper scoring rule. For the application of loss-based Bayesian inference, in which more general forms of loss functions may drive the Bayesian update, we refer the reader to certain of the other literature cited above, namely [Bissiri et al. \(2016\)](#), [Holmes and Walker \(2017\)](#), [Lyddon et al. \(2019\)](#), and [Syring and Martin \(2019\)](#).

### 3.3.3. Bayesian predictive combinations: Beyond BMA

The predictive distributions within the ‘plausible class’ referenced above may characterize a single dynamic structure depending on a vector of unknown parameters,  $\theta$ , or it may constitute weighted combinations of predictives from distinct models, in which case  $\theta$  comprises both the model-specific parameters and the weights. As such, FBP provides a coherent Bayesian method for estimating weighted combinations of predictives via predictive accuracy criteria, and without the need to assume that the true model is spanned by the set of constituent predictives – an assumption that underpins BMA, as we have noted.

A similar motivation underlies other contributions to the extensive Bayesian literature on estimating combinations of predictives that has now developed (and which rivals the large frequentist literature on forecast combinations that has also evolved),<sup>9</sup> with predictive performance – quantified by a range of user-specified measures of predictive accuracy – driving the posterior updating of the weights. Indeed, the Bayesian literature, having access as it does to powerful computational tools, has been able to invoke more complex weighting schemes than can be tackled via frequentist (optimization) methods. Notable contributions, including some also driven by the criterion of predictive calibration ([Dawid, 1982](#); [Dawid, 1985](#); [Gneiting, Balabdaoui, & Raftery, 2007](#)), include [Billio, Casarin, Ravazzolo, and van Dijk \(2013\)](#), [Casarin, Grassi, Ravazzolo, and van Dijk \(2015\)](#), [Casarin, Leisen, Molina, and ter Horst \(2015\)](#), [Casarin, Mantoan, and Ravazzolo \(2016\)](#), [Pettenuzzo and Ravazzolo \(2016\)](#), [Aastveit, Ravazzolo, and van Dijk \(2018\)](#), [Bassetti, Casarin, and Ravazzolo \(2018\)](#), [Baştürk, Borowska, Grassi, Hoogerheide, and van Dijk \(2019\)](#), and [Casarin, Grassi, Ravazzolo, and van Dijk \(2023\)](#). Once again adopting the language of [Bernardo and Smith \(1994\)](#), this literature seeks to move Bayesian predictive combinations beyond the  $\mathcal{M}$ -closed world of BMA to the  $\mathcal{M}$ -open world that accords with the reality of misspecification.

We complete this section by highlighting one particular generalization of BMA that aims not so much to

cater to the  $\mathcal{M}$ -open world but, rather, to remove the fixed-weight restriction that is inherent to BMA. Certain of the references cited above either explicitly allow for the weights attached to the constituent forecasts to evolve over the time period,  $t = 1, 2, \dots, T$ , on which the predictive distribution for  $y_{T+1}$  conditions (e.g. [Billio et al., 2013](#) and [Casarin et al., 2023](#)) or implicitly allow for such a possibility (e.g. [Loaiza-Maya et al., 2021](#)). However, so-called dynamic model averaging (DMA) accommodates time-varying weights via a more direct generalization of BMA, and nests BMA when appropriate settings are activated (see [Koop & Korobilis, 2012](#), page 875 for an illustration of this). We refer the reader to [Raftery, Miroslav, and Ettler \(2010\)](#) for the initial proposal of DMA, to [Koop and Korobilis \(2012\)](#) for the application of the method to forecasting inflation, and to [Nonejad \(2021\)](#) for a recent review of the methodology, with a focus on applications in economics and finance.<sup>10</sup>

### 3.3.4. Bayesian predictive decision synthesis

A third approach that seeks to produce Bayesian predictions without relying explicitly on correct model specification is Bayesian predictive synthesis (BPS) ([Johnson, 2017](#); [McAlinn & West, 2019](#); [McAlinn, Aastveit, Nakajima, & West, 2020](#); [Aastveit, Cross, & van Dijk, 2023](#)), recently expanded to Bayesian predictive decision synthesis (BPDS) by [Tallman and West \(2022\)](#). In particular, BPDS provides a sound decision-theoretic framework for constructing forecast combinations, and can be shown to encompass several commonly suggested Bayesian forecasting approaches.

The starting point of BPDS is the production of a prior distribution over the  $m$ -dimensional unknown outcome  $\mathbf{y}$  – implicitly indexed by  $T + 1$  in a time series forecasting application – and the information set  $\mathcal{H}$ , encoded via the  $J$  predictive models  $\{h_j(\mathbf{y}|\mathbf{x}_j) : 1 \leq j \leq J\}$ , where  $\mathbf{x} = (\mathbf{x}_1, \dots, \mathbf{x}_J)$  denotes the collection of vectors of (possibly latent) dummy variables associated with a decision. The decision maker then constructs a predictive by integrating out  $\mathbf{x}$  using a synthesis function  $\alpha(\mathbf{y}|\mathbf{x})$ :

$$p(\mathbf{y}|\mathcal{H}) = \int_{\mathcal{X}} \alpha(\mathbf{y}|\mathbf{x}) \prod_{j=1}^J h_j(\mathbf{y}|\mathbf{x}_j) d\mathbf{x}_1 \dots d\mathbf{x}_J.$$

The choice of the synthesis function  $\alpha(\cdot|\mathbf{x})$  can be used to drive the analysis. For instance, in the case of forecast combinations, we can take  $h_j(\mathbf{y}|\mathbf{x}) = p_j(\mathbf{y}|\mathbf{x}, \mathcal{M}_j)$ , for some model  $\mathcal{M}_j$ , and then any set of synthesis functions

<sup>9</sup> See [Hall and Mitchell \(2007\)](#), [Ranjan and Gneiting \(2010\)](#), [Geweke and Amisano \(2011\)](#), and [Gneiting and Ranjan \(2013\)](#) for early contributions to the frequentist forecast combination literature, and [Wang, Hyndman, Li, and Kang \(2022\)](#) for a recent review. We note that while the study by [Geweke and Amisano \(2011\)](#) is not explicitly Bayesian, in terms of estimating the optimal predictive combination, it provides important insights into the connection between the ‘optimal linear pool’ and BMA, and also uses Bayesian numerical methods in the production of some of the constituent forecast distributions.

<sup>10</sup> We also refer the reader to [Green \(1995\)](#), [Madigan and Raftery \(1995\)](#), and [George \(2000\)](#) for alternative approaches to catering to model uncertainty in the Bayesian framework. In brief, such approaches (in one way or another) design MCMC samplers to tackle an augmented space in which model uncertainty is incorporated. As a consequence, the computation of any expectation of interest, including that which defines a predictive distribution, automatically factors in all uncertainty associated with both the parameters of each model and the model structure itself. See [Green \(2003\)](#), [Marin, Mengersen, and Robert \(2005\)](#), [Chib \(2011\)](#), and [Fan and Sisson \(2011\)](#) for reviews and more complete referencing.

$\alpha_j(\cdot|\mathbf{x})$  such that the combination density

$$p(\mathbf{y}|\mathcal{H}) = \int_{\mathcal{X}} \frac{\sum_{j=1}^J \omega_j \alpha_j(\mathbf{y}, \mathbf{x}_j | \mathbf{x}) p_j(\mathbf{y} | \mathbf{x}, \mathcal{M}_j)}{\sum_{k=1}^J \omega_k \alpha_k(\mathbf{y}, \mathbf{x}_k | \mathbf{x})} d\mathbf{x}_1 \dots d\mathbf{x}_J.$$

is a valid density for the given weights  $0 < \omega_j < 1$ ,  $\sum_{j=1}^J \omega_j = 1$ . Specific choices of  $\alpha_j(\mathbf{y}, \mathbf{x}_j | \mathbf{x})$  then produce different forecast combination methods (see [Johnson, 2017](#) for a discussion). For example, in the case of [McAlinn and West \(2019\)](#) and [McAlinn et al. \(2020\)](#), the synthesis function is taken to be the density of a (possibly multivariate) dynamic linear factor model.

In an attempt to ‘focus’ the BPDS approach towards decisions that are tailored to a specific user-chosen loss function underlying the analysis or decision at hand, [Tallman and West \(2022\)](#) propose taking as their synthesis function  $\alpha_j(\mathbf{y}, \mathbf{x}_j | \mathbf{x}) = \exp\{\tau'(\mathbf{x})S_j(\mathbf{y}, \mathbf{x}_j)\}$ , where the score  $S_j(\mathbf{y}, \mathbf{x}_j)$  is a  $k$ -dimensional vector that measures the utility one receives from realizing outcome  $\mathbf{y}$  under decision  $\mathbf{x}_j$ , and  $\tau(\mathbf{x})$  is a vector that weights the directional relevance of  $S_j(\mathbf{y}, \mathbf{x}_j)$ .

While the BPS framework, as a whole, can set the tenor of the predictions towards dynamic forecast updates that produce predictions tailored to a loss function of interest, via the choice of synthesis function  $\alpha(\cdot|\cdot)$ , BPS is ultimately tied to a likelihood-type framework, or at least a log-loss function, due to the presence of the latent variables  $\mathbf{x}$ , which must be integrated out via assumed predictive models,  $p_j(\mathbf{y}|\mathbf{x}, \mathcal{M}_j)$ , and with these individual predictives produced using likelihood-based Bayesian methods. While the BPDS approach can somewhat circumvent the reliance on the likelihood, due to its ability to focus on specific scores, this approach appears to be distinct from methods that entirely replace the likelihood function in the update. Therefore, a very interesting research path would involve combining the methods based on generalized posteriors discussed in Section 3.3.2 with the BPS framework.

#### 4. Selective discipline-specific reviews of Bayesian forecasting

Having established the necessary details regarding the production of Bayesian forecasts in general contexts, we now review how this general probabilistic mechanism is employed to produce Bayesian forecasts in several important empirical fields. In order to produce a comprehensive and up-to-date review of each area, a range of discipline experts have been invited to write the various sections, with the authorship flagged in the section headings. This means that the style of coverage differs somewhat across sections, as suits the topic and as fits with the perspective of the authors. However, we have aimed to retain notation that (as far as possible) is both consistent across sections and consistent with the notation used in the earlier parts of the paper and in the technical appendix, plus to ensure that the basic layout of all sections is the same. As noted above – other than in Section 4.3, in which cross-sectional consumer choice data are modeled; and in Section 4.4, in which spatial models are briefly referenced – time series problems and forecasting are the primary focus.

#### 4.1. Macroeconomics (Florian Huber and Gary Koop)

Central banks and other policy institutions routinely collect vast amounts of time series data on key macroeconomic outcomes. One stylized fact is that these data sets often display substantial co-movements and this calls for modeling all these series jointly to produce accurate point and density forecasts. This, however, leads to large-scale models that are prone to overfitting, ultimately resulting in weak out-of-sample forecasting performance. This helps explain the popularity of Bayesian methods for macroeconomic forecasting. They can easily handle many parameters and, through appropriate prior choice, deal effectively with questions related to model and specification uncertainty in macroeconomic settings.

At a high level of generality, there are two modeling approaches used by macroeconomic forecasters. The first uses reduced-form models and imposes relatively little economic structure on the data. The second uses structural models such as dynamic stochastic general equilibrium (DSGE) models that are often estimated through Bayesian techniques; see, among many others, [Adolfson, Lindé, and Villani \(2007\)](#), [Smets and Wouters \(2007\)](#) and [Del Negro, Hasegawa, and Schorfheide \(2016\)](#). However, reduced-form approaches have proved more popular and, in this section, our focus will be on them.

As stated above, macroeconomists are typically interested in modeling the joint evolution of a set of macroeconomic quantities. To set up a general framework for understanding the types of models used for forecasting, assume that an  $M$ -dimensional vector  $\mathbf{y}_t$  is related to a  $K$ -dimensional vector of explanatory variables  $\mathbf{x}_t$  through

$$\mathbf{y}_t = g(\mathbf{x}_t) + \boldsymbol{\varepsilon}_t, \quad (16)$$

where  $g: \mathbb{R}^K \rightarrow \mathbb{R}^M$  is a function, and  $\boldsymbol{\varepsilon}_t$  is  $\mathcal{N}(\mathbf{0}_M, \boldsymbol{\Sigma}_t)$ .<sup>11</sup> This general specification nests most important reduced-form models commonly used in macroeconomics and can be used to explain the main issues that arise.

For instance, if  $\mathbf{x}_t = (\mathbf{y}'_{t-1}, \dots, \mathbf{y}'_{t-p})'$  contains  $p$  lags of  $\mathbf{y}_t$ ,  $g(\mathbf{x}_t) = \mathbf{A}\mathbf{x}_t$  is a linear function with  $M \times K (= Mp)$  coefficient matrix  $\mathbf{A}$ , and  $\boldsymbol{\Sigma}_t = \boldsymbol{\Sigma}$  is constant over time, we have a standard vector autoregressive (VAR) model. If we set  $\mathbf{x}_t = \mathbf{f}_t$ , with  $\mathbf{f}_t$  denoting a set of  $Q \ll M$  latent factors and  $g(\mathbf{f}_t) = \mathbf{A}\mathbf{f}_t$  is linear with  $\mathbf{A}$  being an  $M \times Q$  matrix of factor loadings and  $\mathbf{f}_t$  evolves according to some stochastic process (such as a VAR), we end up with a dynamic factor model (DFM, see [Stock & Watson, 2011](#)). Factor-augmented VARs ([Bernanke, Boivin, & Elias, 2005](#)) combine a VAR with a DFM. The dependent variables in the VAR part of the model are a subset of  $\mathbf{y}_t$  plus a small number of factors.

Traditionally, VARs and factor models have been linear and homoskedastic. But there is a great deal of empirical evidence in most macroeconomic data sets of parameter change, both in the conditional mean and the conditional variance. This can be accommodated through particular

<sup>11</sup> Note that the Gaussianity assumption is not essential; mixtures of Gaussian distributions, for example, can be used to produce flexible error distributions if deemed necessary (see, for example, [Clark, Huber, Koop, Mercellino, & Pfarrhofer, 2022b](#) and [Lenza & Primiceri, 2022](#)).

choices for  $g$  and  $\Sigma_t$ . For the latter, stochastic volatility processes have proved particularly popular. For the former, various parametric forms for  $g$  lead to time-varying parameter VARs (TVP-VARs), which assume that the coefficients of the VAR evolve according to a random walk. But it is also worth noting that there is an increasing amount of literature that assumes that  $g$  is unknown and uses Bayesian non-parametric methods to uncover its form (see, for example, Kalli & Griffin, 2018, Adrian, Boyarchenko, & Giannone, 2021, and Huber, Koop, Onorante, Pfarrhofer, & Schreiner, 2023).

If we set  $M = 1$ , we obtain single-equation time series regressions, which are particularly popular in inflation forecasting (e.g. based on the Phillips curve). If we additionally set  $x_t = 1$  and allow for time-varying parameters, we can obtain models such as the unobserved components stochastic volatility (UCSV) model of Stock and Watson (2007), which is commonly used to forecast inflation (for recent applications, see Chan, Koop, & Potter, 2013, Stock & Watson, 2016, and Huber & Pfarrhofer, 2021).

This general framework defines a class of likelihood functions. As per the outline in Section 2.1, Bayesian forecasting involves multiplying a chosen likelihood function by an appropriate prior to produce a posterior which can be used to produce the predictive density. The choice of prior and computational method used for posterior and predictive inference will be case-specific and we will have more to say about some interesting cases below. But a few general comments are worth noting here. First, the choice of prior matters much more in models, such as the large VAR, that have a large number of parameters relative to the number observations, than in models with fewer parameters, such as the UCSV model or the DFM. Second, for linear homoskedastic models with conjugate priors, analytical formulae for the posterior and the one-step-ahead predictive density are available. For all other cases, MCMC methods are available. These take the general form outlined in Section 3.1. However, as noted in Section 3.2, MCMC methods typically do not scale well and can be computationally slow in models involving large numbers of parameters (such as large VARs) or large numbers of latent states (such as TVP-VARs). Thus, the focus of many recent papers has been on developing either improved MCMC algorithms or approximate VB methods for speeding up computation. Third, our discussion so far focuses on forecasting with a single model. In practice, it is common to find that forecasts improve if many models are combined. Thus, either BMA or, alternatively, the methods outlined in Section 3.3.3 are commonly used by macroeconomic forecasters.

With this general framework established, it is worthwhile to offer some additional detail about some of the most important 21st century developments and a discussion of how they have led to improvements in macroeconomic forecasting.

**Large VARs.** Going back to early work such as Doan, Litterman, and Sims (1984), Bayesian VARs have been used successfully in a variety of macroeconomic forecasting applications. Recently, they have enjoyed even greater

popularity, due to the rise of the large VAR. The pioneering large VAR paper was Bańbura, Giannone, and Reichlin (2010). Subsequently, dozens of papers have used large VARs for macroeconomic forecasting (see, among many others, Carriero, Kapetanios, & Marcellino, 2009, Koop, 2013b, Carriero, Clark, & Marcellino, 2015, Giannone, Lenza, & Primiceri, 2015, and Hauzenberger, Huber, & Onorante, 2021). Large VARs, involving dozens or even hundreds of dependent variables, have been found to forecast well and improve upon single-equation techniques and DFMs. Large VARs are heavily over-parameterized and, thus, Bayesian prior shrinkage has been essential in ensuring their forecasting success. We discuss priors below, but at this point we highlight the fact that the use of large Bayesian VARs has been one of the major recent developments in macroeconomic forecasting.

**Prior shrinkage in VARs.** Many different priors have been used with VARs. Traditionally, natural conjugate priors in the Minnesota tradition were used, since these allowed for analytical posteriors and one-step-ahead predictives. Definitions of these priors and discussions of their properties are available in standard sources such as Koop and Korobilis (2010) and Dieppe, van Roye, and Legrand (2016). These priors are subjective and require the user to select prior hyperparameters, most importantly those relating to the strength of prior shrinkage. In recent years, a range of alternative priors have been proposed that are more automatic, requiring fewer subjective prior choices by the researcher. For instance, Giannone et al. (2015) develop methods for estimating shrinkage parameters in conjugate priors, thus avoiding the need for their subjective elicitation. Chan (2022) also uses a conjugate prior and develops methods for selecting shrinkage parameters using a prior which relaxes some of the restrictive assumptions of the Minnesota prior. There are also a range of methods which automatically decide on the optimal degree of shrinkage for each VAR coefficient. These are the global-local shrinkage priors which are widely used with regressions and in machine learning applications, and increasingly used with VARs.<sup>12</sup> Global-local shrinkage priors have the form

$$a_j \sim \mathcal{N}(0, \psi_j \lambda), \quad \psi_j \sim f_1, \quad \lambda \sim f_2,$$

where  $a_j$  is the  $j$ th VAR coefficient,  $\lambda$  controls global shrinkage since it is common to all coefficients, and  $\psi_j$  controls local shrinkage since it is specific to the  $j$ th coefficient. The densities  $f_1$  and  $f_2$  are mixing densities, and a large range of choices of them have been proposed. One choice leads to stochastic search variable selection, used with VARs in George, Sun, and Ni (2008), Koop (2013a), and Korobilis (2013), and many other references. Other choices lead to the Dirichlet-Laplace prior used with VARs by Kastner and Huber (2021), or the normal-gamma and horseshoe priors used in Huber and Feldkircher (2019) and Cross, Hou, and Poon (2020); and there are many others. Since these priors are Gaussian at the first layer of

<sup>12</sup> They are also used with DFMs to select the number of factors.



the hierarchy, textbook MCMC algorithms for all the VAR parameters can be implemented easily.<sup>13</sup>

*Adding stochastic volatility (SV).* The other main development that has had a tremendous impact on applied macroeconomic forecasting in the 21st century is the development of models such as VARs that incorporate parameter change and non-linearity. Put simply, the macroeconomic world is rarely linear and homoskedastic, and models that relax these assumptions have been found to improve macroeconomic forecasting. These improvements lie not only in point forecasts, but more importantly in density forecasts. Given the increasing interest, by central banks and academics alike, in issues such as forecast uncertainty and tail risk, the fact that these new models produce more accurate predictive densities increases their value.

A popular specification for VARs with SV involves factorizing the error variance–covariance matrix as  $\Sigma_t = \mathbf{A}_0 \mathbf{H}_t \mathbf{A}_0'$ , with  $\mathbf{A}_0$  being a lower triangular matrix with unit diagonals<sup>14</sup> and  $\mathbf{H}_t = \text{diag}(e^{h_{1t}}, \dots, e^{h_{Mt}})$  being a diagonal matrix with log-volatilities evolving according to simple stochastic processes such as independent random walks or AR(1) processes. In an important contribution, Clark (2011) considers a VAR-SV and finds it to produce accurate point and density forecasts relative to homoskedastic models, with gains being particularly pronounced using forecast metrics involving the entire predictive density. Building on this insight, several other researchers have analyzed the role of heteroskedasticity in macroeconomic forecasting in VARs (see, for example, Clark & Ravazzolo, 2015 and Chiu, Mumtaz, & Pinter, 2017) and confirm the result that using SV pays off when the focus is on obtaining accurate density forecasts. However, a problem with the standard SV specification is that the computational burden relative to homoskedastic VARs is increased enormously. This makes it difficult to do Bayesian forecasting with large VARs with SV. As a remedy, Carriero, Clark, and Marcellino (2016) propose a simple common stochastic volatility (CSV) specification that assumes the shock variances to be driven by a single common volatility factor, maintaining conjugacy and thus leading to computationally efficient MCMC algorithms. They acknowledge that this model is simplistic but show that it yields much more accurate forecasts than homoskedastic VARs in a standard US macroeconomic forecasting application.

To gain more flexibility, researchers have developed algorithms that allow for estimating large VARs with  $M$  independent SV processes. Carriero, Clark, and Marcellino (2019) propose techniques that permit equation-by-equation estimation of such VARs and thus render estimation of larger models with SV feasible. Modified versions of this algorithm form the basis of several recent papers that combine large data sets with SV for macroeconomic forecasting (see, among others, Huber &

Feldkircher, 2019, Chan, 2021, and Chan, Koop, & Yu, 2023).

*Adding time variation in the VAR coefficients.* The previous discussion emphasized that capturing changing error variances is key for obtaining precise forecasts. However, it may also be important to allow for structural change in the VAR coefficients themselves. One popular multivariate model that captures both changes in the VAR coefficients and error variances is the TVP-VAR-SV model proposed by Primiceri (2005), which assumes that the VAR coefficients  $\beta_t = \text{vec}(\mathbf{A}_t)$  are time-varying and evolve according to a multivariate random walk, while  $\Sigma_t$  is a multivariate SV process. This model is a multivariate state-space model which can be estimated using adaptations of the techniques outlined in Section 3.1.3. The innovations to the states govern the amount of time variation in the parameters. Various shrinkage priors (often based on the global–local shrinkage priors discussed above) have been proposed that allow for a data-based decision as to whether time variation in a corresponding coefficient is necessary or not. These priors are typically elicited on the non-centered parameterization of the state-space model (see Frühwirth-Schnatter & Wagner, 2010) and can help minimize overfitting concerns and produce improved forecasts.

D'Agostino, Gambetti, and Giannone (2013) made an important early contribution to the macroeconomic forecasting literature using TVP models. They use a small TVP-VAR with SV and show that it produces more accurate point predictions, outperforming simpler univariate benchmarks and constant-parameter VARs. One key shortcoming of this model, however, is that it only uses a small information set. This has led to several researchers proposing new methods that can be used in higher dimensions. Various approaches are possible, including models that restrict the TVP process (e.g. by imposing a factor structure, which allows for time variation in a large number of parameters to be driven by a low number of factors; see Chan, Eisenstat, & Strachan, 2020). As mentioned above, shrinkage priors are used to keep the curse of dimensionality in check. These priors are typically used after transforming the model to allow for equation-by-equation estimation. Such approaches mean that fairly high-dimensional TVP-VARs can be estimated without risk of overfitting, and in a reasonable amount of time. MCMC-based forecasting with large TVP-VARs and regressions is also an active field of research, and different shrinkage methods and advances in computation have led to improvements in the forecasting performance of TVP models (see, among many others, Huber, Koop, & Onorante, 2021 and Hauenberger, Huber, Koop, & Onorante, 2022). However, it is worth noting that if computation does become a concern, approximate methods (e.g. using the VB methods outlined in Section 3.2.3) can be used. Approaches that avoid the need for MCMC are developed by Koop and Korobilis (2013, 2023). In the former paper the authors propose large approximate TVP-VARs based on forgetting factors, whereas in the latter they use VB

<sup>13</sup> In large VARs with global–local shrinkage priors, MCMC methods can nevertheless be very slow, with much faster VB methods developed in Gefang, Koop, and Poon (2022).

<sup>14</sup>  $\mathbf{A}_0$  can also be time-varying.

techniques to forecast inflation with large TVP regression models.

**Bayesian non-parametric VARs.** Up to this point, we have assumed that the conditional mean function  $g$  takes a known form. However, it could be that the functional form is unknown. Bayesian non-parametric techniques, such as Bayesian additive regression trees (BART; see [Chipman, George, & McCulloch, 2010](#)), Gaussian processes and kernel regressions ([Adrian et al., 2021](#)), or infinite mixtures ([Kalli & Griffin, 2018](#)), allow the researcher to uncover such unknown functional forms and produce precise macroeconomic forecasts. In general, they have had great success, but they have been found to be particularly useful in studies that focus on the tails of predictive distributions or on the handling of outliers such as the ones experienced during the pandemic (see, for example, [Clark, Huber, Koop, & Marcellino, 2022a](#) and [Huber et al., 2023](#)).

[Kalli and Griffin \(2018\)](#) propose a non-parametric VAR that builds on an infinite mixture model, with the mixture weights being driven by the lagged endogenous variables. They show, using US and UK data, that their model yields competitive forecasts, with accuracy gains in terms of point and density predictions increasing sharply for higher forecast horizons. [Clark et al. \(2022b\)](#) use BART-based VARs to perform tail forecasting of US output, unemployment and, inflation in real time, finding that non-parametric techniques work well in the tails and for higher-order forecasts. With a particular focus on predictive accuracy during the pandemic, [Huber et al. \(2023\)](#) develop mixed-frequency non-parametric VARs and show that these models yield substantially more precise nowcasts during the Covid-19 period.

**Conclusions and further directions.** We outlined how Bayesian methods have been used successfully for macroeconomic forecasting. Most of the discussion related to VARs, which are a class of models where Bayesian methods have proved particularly popular. But it is worth noting that empirically relevant extensions (e.g. SV or TVP) can be added to other multivariate time series models such as DFMs or FAVARs, as can the VAR prior shrinkage methods (e.g. global-local shrinkage methods) we discussed. It is also worth noting that we focused on models that do not restrict the coefficients. However, restricted VARs are often used for forecasting. For instance, vector error correction models (which impose cointegrating restrictions) or multi-country VARs such as global VARs are restricted VARs.

We also focused on forecasting as opposed to the closely related field of nowcasting. Mixed-frequency VARs, which jointly model quickly released, high-frequency variables (e.g. monthly variables such as surveys, employment, and inflation) and slowly released, low-frequency variables (e.g. quarterly variables such as GDP), have proved very popular with nowcasters. Bayesian methods are typically used with such models (see, for example, [Schorfheide & Song, 2015](#), [Koop, McIntyre, Mitchell, & Poon, 2020](#), [McCracken, Owyang, & Sekhposyan, 2021](#), and [Huber et al., 2023](#)), and in real-time nowcasting exercises they tend to perform well.

#### 4.2. Finance (John Maheu, Worapree Maneesoonthorn, and Gael Martin)

A pertinent question in financial analysis is whether the risks associated with financial assets – and the prices of those risks – are predictable in ways that are useful in applications such as portfolio allocation, risk management, and derivative pricing. With risk factors typically being represented as latent distributional features of observable financial variables, it follows that two key goals in the statistical analysis of financial problems are the accurate prediction of latent distributional features, and the development of complex, non-linear state-space models to underpin this prediction.

Both of these goals lend themselves naturally to a Bayesian treatment, given, in turn, the automatic production of predictive *distributions* via the Bayesian paradigm and the swathe of computational methods available to estimate complex models, most notably those with a latent variable structure. In particular, the growth in financial derivatives markets from the 1990s onwards has generated the need to model the underlying asset as a continuous time process, almost always augmented with a continuous time process for the asset volatility, and often via a jump diffusion. Such models – whilst convenient in the sense of allowing for closed-form solutions for derivative prices – are challenging from a statistical point of view, given that they typically need to be treated as a (discretized) non-linear state-space model, and may require multiple sources of data to enable separate identification of model parameters and risk premia. Estimation of and forecasting with such models is nevertheless computationally feasible via Bayesian methods, with MCMC algorithms of one form or another forming the backbone of the early treatments ([Eraker, 2001](#); [Eraker, Johannes, & Polson, 2003](#); [Eraker, 2004](#); [Forbes, Martin, & Wright, 2007](#); [Johannes, Polson, & Stroud, 2009](#)).

We refer the reader to [Johannes and Polson \(2010\)](#) and [Jacquier and Polson \(2011\)](#) for comprehensive reviews of the application of Bayesian methods in finance up to the first decade of the 21st century. The coverage includes, in short, Bayesian approaches to portfolio allocation, return predictability, asset pricing, volatility, covariance, ‘beta’ and ‘value-at-risk’ prediction, continuous time models (and discretized versions thereof), interest rate modeling, and derivative (e.g. option) pricing. Our goal in the current review is to outline the more recent advances that have evolved over the last decade, in particular those that have exploited (in one way or another) new methodological advances, new sources of data, and modern computational techniques. In order, we shall briefly review the following: the use of diverse data sets, including derivative prices and high-frequency measures of financial quantities; the treatment of DGPs that are unavailable in closed form; the analysis of high-dimensional models; and the application of non-parametric modeling.

**Multiple sources of financial data.** It is now a well-established fact that the constant volatility feature of a geometric diffusion process for a financial asset price

is inconsistent with both the observed dynamics in return volatility and the excess kurtosis and skewness that characterizes the typical empirical return distribution; see [Bollerslev, Chou, and Kroner \(1992\)](#) for an early review. The option pricing literature supports this finding, with certain empirical regularities, such as ‘implied volatility smiles’, seen as evidence that asset prices deviate from the geometric Brownian motion assumption that underlies the [Black and Scholes \(1973\)](#) option price ([Bakshi, Cao, & Chen, 1997](#); [Hafner & Herwartz, 2001](#); [Lim, Martin, & Martin, 2005](#)). Hence, the 21st century has seen the proliferation of many alternative specifications for asset prices and associated theoretical derivative prices, most of which are nested in a general framework of (discretized) bivariate jump diffusion models for the asset itself and its volatility. Allied with these developments has been the growth in access to transaction-level high-frequency data – in both the spot and options markets – which, in itself, has spawned new approaches to inference and forecasting in the financial sphere.

The Bayesian literature has brought to bear on this problem the power of computational methods – both established, and more recent – to enable the multivariate state-space models that have emerged from this literature to be estimated, and probabilistic predictions of all dynamic variables – the return itself, volatility, random jumps (in either the return or the volatility, or both), and various risk premia – to be produced. With reference to the generic notation for a state-space model in (7) and (8), Bayesian approaches over the last decade can be categorized according to the specification adopted for the (multivariate) measurement at time  $t$ ,  $\mathbf{y}_t$ , and, hence, for the (multivariate) state,  $\mathbf{z}_t$ , being modeled and forecast. Some work exploits data from both the spot and options market to predict volatility and its risk premia ([Maneesoonthorn, Martin, Forbes, & Grose, 2012](#)), and option prices ([Yu, Li, & Wells, 2011](#); [Carverhill & Luo, 2023](#)<sup>15</sup>). Other work combines low-frequency daily observations on returns with high-frequency measures of volatility and/or price jumps to predict (in some combination) returns, volatility, and the size and occurrence of price jumps ([Jin & Maheu, 2013](#); [Maneesoonthorn, Forbes, & Martin, 2017](#); [Frazier, Maneesoonthorn, Martin, & McCabe, 2019](#)). Further work combines daily returns with futures prices in predicting various financial quantities of interest ([Fileccia & Sgarra, 2018](#); [Gonzato & Sgarra, 2021](#)).

*Financial models that are ‘unavailable’.* All but one of the papers cited in the previous paragraphs share a common feature – namely, a DGP that can be expressed as a probability density (or mass) function. With reference to (9), it is the availability of a closed form for  $p(\mathbf{y}_{1:T}, \mathbf{z}_{1:T}|\theta) = p(\mathbf{y}_{1:T}|\mathbf{z}_{1:T}, \theta)p(\mathbf{z}_{1:T}|\theta)$  that renders feasible the MCMC methods used in the said works. In contrast, [Frazier et al. \(2019\)](#) adopt a process for the latent

log-volatility that is driven by an  $\alpha$ -stable innovation, such that  $p(\mathbf{z}_{1:T}|\theta)$  is unavailable and MCMC infeasible as a consequence. Instead, ABC is adopted for inference, and an approximate predictive of the form of (12) is produced instead. In addition to providing theoretical validation of the approach, the authors demonstrate, in range of different simulation settings, that despite inaccuracy at the posterior level, the approximate predictive is always a very close match to the exact predictive. Related work in which an ABC method is used to conduct forecasting appears in [Canale and Ruggiero \(2016\)](#), [Kon Kam King, Canale, and Ruggiero \(2019\)](#), [Virbickaitė, Ausín, and Galeano \(2020\)](#), and [Pesonen et al. \(2022\)](#). ABC treatment of a conditional likelihood for a time series of financial returns,  $p(\mathbf{y}_{1:T}|\mathbf{z}_{1:T}, \theta)$ , that is unavailable in closed form is also investigated by [Creel and Kristensen \(2015\)](#), [Martin et al. \(2019\)](#), and [Chakraborty, Nott, Drovandi, Frazier, and Sisson \(2022\)](#), with [Chakraborty et al. \(2022\)](#) proposing a modularized version of ABC. For other recent Bayesian treatments of intractable models of this sort that continue to exploit MCMC principles (with or without an ABC component), see [Vankov, Guindani, and Ensor \(2019\)](#) and [Müller and Uhl \(2021\)](#).<sup>16</sup>

*Large financial models.* Thus far, we have reviewed Bayesian treatments of models for single financial assets. That is, the models may have specified multiple latent components, and potentially multiple measurements, but they still aim to explain (and forecast) quantities related to a *single* asset. Models for multiple assets are also critically important in financial applications, with the relationship between financial assets determining the extent to which diversification can be achieved, as well as how risks permeate the various sectors of the financial market. Indeed, Bayesian methods are particularly suitable for dealing with such multivariate models, since the dimensionality of  $\mathbf{z}_{1:T}$  is typically much larger than that of  $\mathbf{y}_{1:T}$  and, hence, challenging to deal with via any other means.

[Chib, Omori, and Asai \(2009\)](#) provide an early review of the Bayesian analysis of multivariate SV models, with all work up to this point utilizing traditional MCMC techniques, and the statistical and predictive analysis limited to relatively low-dimensional systems (up to ten assets). Subsequent work has focused on the development of more flexible multivariate distributions ([Nakajima, 2017](#)) and the use of sparse factor structures and shrinkage priors in constructing larger-dimensional models ([Zhou, Nakajima, & West, 2014](#); [Kastner, Frühwirth-Schnatter, & Lopes, 2017](#); [Baştürk, Borowska, Grassi, Hoogerheide, & van Dijk, 2019](#)). More recently, with the advances made in VB methods, inference and prediction in very large-dimensional financial models is now possible ([Gunawan, Kohn, & Nott, 2021](#); [Chan & Yu, 2022](#); [Frazier, Loaiza-Maya, & Martin, 2022](#); [Quiroz, Nott, & Kohn, 2022](#); [Zhang, Smith, Maneesoonthorn, & Loaiza-Maya, 2023](#)). There is

<sup>15</sup> We note that whilst a time series model is constructed in the case of these two references, the (out-of-sample) prediction of option prices is across the cross-section of strike prices and maturities. We also make note of [Fulop and Li \(2019\)](#), who exploit spot and options data to produce filtered estimates (as opposed to strictly out-of-sample predictions) of latent volatility and price jump intensity.

<sup>16</sup> The citation of [Creel and Kristensen \(2015\)](#), [Martin et al. \(2019\)](#), [Vankov et al. \(2019\)](#), and [Müller and Uhl \(2021\)](#) is relevant to this review, despite these references not having an explicit component on forecasting.



also a growing interest in the prediction of co-movements of various sorts – with [Bernardi, Gayraud, and Petrella \(2015\)](#) predicting the interdependence between US stocks with Bayesian time-varying quantile regressions, [Geraci and Gnabo \(2018\)](#) capturing and predicting the interconnectedness of financial institutions through Bayesian time-varying VARs, and [Alexopoulos, Dellaportas, and Paspaliopoulos \(2022\)](#) modeling and predicting common jump factors in a large panel of financial returns.

*Bayesian non-parametric modeling in finance*. As noted, simple parametric assumptions such as additive Gaussian innovations are inconsistent with the stylized features of financial data. Although more suitable non-Gaussian/non-linear models can be built (as highlighted above), Bayesian nonparametric modeling allows for further flexibility via the incorporation of Dirichlet process mixture (DPM) structures. Such an approach has been shown to provide robustness to distributional assumptions and can improve point forecasts, but the main gain has been significant improvements in the accuracy of predictive densities, and of risk measures derived from those densities. The advancement of the literature in this direction has been aided by the stick-breaking representation ([Sethuraman, 1994](#)) and the introduction of the slice sampler ([Walker, 2007](#); [Kalli, Griffin, & Walker, 2011](#)).

[Jensen and Maheu \(2010\)](#) introduce an extension to a standard SV model to capture the unknown return innovation distribution via a DPM. The DPM specification has also been inserted into other popular models in finance – with [Jensen and Maheu \(2014\)](#) adopting a DPM to jointly model the return and future log-volatility distribution, [Delatola and Griffin \(2013\)](#) capturing the so-called leverage effect, [Ausín, Galeano, and Ghosh \(2014\)](#) applying a DPM to univariate GARCH models, and [Kalli and Griffin \(2015\)](#) using Bayesian non-parametric modeling to aggregate autoregressive processes to produce an SV model with long-range dependence. Extensions to multivariate financial models have also occurred: in a multivariate GARCH setting in [Jensen and Maheu \(2013\)](#), and in a Cholesky-type multivariate SV model in [Zaharieva, Trede, and Wilfling \(2020\)](#).

A potential drawback of the DPM model is that it neglects time dependence in the unknown distribution. An important extension of the DPM prior is the hierarchical Dirichlet process of [Teh, Jordan, Beal, and Blei \(2006\)](#), which allows for the construction of a prior for an infinite hidden Markov model (IHMM) for time dependence in a flexible manner. The introduction of the beam sampler of [Van Gael, Saatchi, Teh, and Ghahramani \(2008\)](#), which extends the slice sampler, renders conventional posterior sampling methods for finite-state Markov switching models ([Chib, 1996](#)) feasible in the IHMM. The IHMM structure has been used to model the univariate GARCH distribution ([Dufays, 2016](#)) and the multivariate GARCH distribution ([Li, 2022](#)), and to provide a non-parametric model for realized measures, including realized covariance matrices ([Jin & Maheu, 2016](#); [Liu & Maheu, 2018](#); [Jin, Maheu, & Yang, 2019](#)), with all papers documenting very large improvements in density forecast accuracy from the IHMM. Other applications of the IHMM include: [Shi and](#)

[Song \(2016\)](#), who use the IHMM to date and forecast speculative bubbles, and who also adopt a version with GARCH effects; [Yang \(2019\)](#), who studies the relationship between stock returns and real growth with a multivariate IHMM model; and, more recently, [Jin, Maheu, and Yang \(2022\)](#), who employ the DPM prior in the infinite Markov pooling of predictive distributions, with forecasting applications to interest rates, realized covariances, and asset returns. Other approaches to time dependence in Bayesian non-parametrics for finance include [Griffin and Steel \(2011\)](#), who introduce a time-dependent stick-breaking process in a general setting and develop an SV model for returns. More recently, [Sun, Kim, and Lee \(2020\)](#) use a weighted DPM to forecast return distributions, while [Zamenjani \(2021\)](#) allows for lagged covariates to impact the weights in the DPM model through a probit stick-breaking process.

#### 4.3. Marketing (*Rubén Loaiza-Maya and Didier Nibbering*)

Bayesian methods are applied to a wide range of marketing problems; see [Rossi and Allenby \(2003\)](#) for a review of the early literature. More recently, these methods have been increasingly used for the purpose of prediction, for instance in customer choice behavior ([Toubia, Iyengar, Bunnell, & Lemaire, 2019](#); [Araya, Elberg, Noton, & Schwartz, 2022](#)), customer demand ([Posch, Truden, Hungerländer, & Pilz, 2022](#)), customer satisfaction ([Mittal, Han, Lee, & Sridhar, 2021](#)), dynamic pricing ([Bastani, Simchi-Levi, & Zhu, 2022](#)), advertising effectiveness ([Danaher, Danaher, Smith, & Loaiza-Maya, 2020](#); [Loaiza-Maya, Smith, Nott, & Danaher, 2022](#)), and recommender systems ([Ansari, Li, & Zhang, 2018](#)). Given the large variety of marketing applications, we focus in this section on the modeling of customer choice, to illustrate the key principles of Bayesian prediction in marketing problems.

A common problem in marketing is that of setting the price level of a set of products so that total profits are maximized. To estimate these optimal prices, predictions of how customers will react to price changes are crucial. Predictions of customer choices under different marketing environments can be constructed by choice models. These models are estimated using data about the product choices of customers in the marketplace, a survey, an experiment, etc. ([Rossi, Allenby, & McCulloch, 2012](#)).

An example of a prediction of interest in this context is the predicted purchase probability of a customer for a particular product as a function of its own price or the price of another product. The predicted purchase probability can be constructed for a customer for which only a few choices are observed, or for a new customer for which we do not observe choices in the data.<sup>17</sup>

The two most popular models used to predict choice behavior are the multinomial logit and multinomial probit

<sup>17</sup> Although this section, as noted in the introduction, focuses on prediction using cross-sectional data, choice models can also be applied to the forecasting of future choice probabilities by using time series data ([McCormick, Raftery, Madigan, & Burd, 2012](#)) or panel data ([Gilbride & Allenby, 2004](#); [Terui, Ban, & Allenby, 2011](#)).



models. The multinomial logit model imposes the independence of irrelevant alternatives (IIA) property (McFadden, 1989), which means that it cannot capture general substitution patterns among choice alternatives. The IIA property of this model can be relaxed under certain assumptions by extending the multinomial logit model to a nested logit model (Poirier, 1996; Lahiri & Gao, 2002) or a random parameter logit model (Train, 2009).

On the other hand, the multinomial probit model does not impose the IIA property, and as such is commonly used in the analysis of economic choice behavior, where complementary and substitution effects are important. For instance, the multinomial probit model has recently been used in the analysis of car choices (Karmakar, Kwon, Mukherjee, & Siddarth, 2021), grocery brand choices (Miyazaki, Hoshino, & Böckenholt, 2021), employment choices (Mishkin, 2021), and car parking choices (Paleti, 2018). The remainder of this section presents a review of Bayesian prediction based on the multinomial probit model.

**Multinomial probit model specification.** The variable of interest is  $y_i \in \{0, 1, 2, \dots, J\}$ , which indicates the choice made by individual  $i$  among a set of  $J + 1$  alternatives. This choice is modeled to be conditional on a set of  $J$  latent utilities  $\mathbf{z}_i = (z_{i1}, \dots, z_{iJ})'$ , so that the conditional pmf is defined as

$$p(y_i | \mathbf{z}_i) = \begin{cases} I[z_{iy_i} = \max(\mathbf{z}_i)] & \text{if } \max(\mathbf{z}_i) > 0, \\ I[y_i = 0] & \text{if } \max(\mathbf{z}_i) \leq 0, \end{cases} \quad (17)$$

where  $p(y_i | \mathbf{z}_i) = \Pr(Y_i = y_i | \mathbf{z}_i)$ ,  $z_{iy_i}$  is the  $y_i$ -th element of  $\mathbf{z}_i$ , with  $y_i > 0$ , and  $I[A]$  is one if statement  $A$  is true and zero otherwise. The base category  $j = 0$  is one of the choice alternatives, which is selected *a priori*. The base category is observed whenever all the latent utilities are less than zero.

The utilities are expressed in terms of  $r$  predictors via a linear Gaussian model,

$$p(\mathbf{z}_i | X_i, \boldsymbol{\theta}) = \phi_j(\mathbf{z}_i; X_i \boldsymbol{\beta}, \boldsymbol{\Sigma}), \quad (18)$$

where  $\phi_j(\mathbf{z}; \boldsymbol{\mu}, C)$  denotes a  $J$ -variate normal density with mean  $\boldsymbol{\mu}$  and covariance matrix  $C$ ,  $X_i$  is a  $J \times r$  matrix of predictor values,  $\boldsymbol{\beta}$  is an  $r$ -dimensional vector of coefficients, and  $\boldsymbol{\Sigma}$  is a covariance matrix that captures complementary and substitution effects between the choice alternatives.

Combined, (17) and (18) give rise to the augmented likelihood function of the multinomial probit model:

$$p(\mathbf{y}, \mathbf{z} | \boldsymbol{\theta}, X) = \prod_{i=1}^n p(y_i | \mathbf{z}_i) p(\mathbf{z}_i | X_i, \boldsymbol{\theta}), \quad (19)$$

where  $\boldsymbol{\theta} = \{\boldsymbol{\beta}, \boldsymbol{\Sigma}\}$ ,  $\mathbf{y} = \{y_i\}_{i=1}^n$ ,  $\mathbf{z} = \{\mathbf{z}_i\}_{i=1}^n$ , and  $X = \{X_i\}_{i=1}^n$ , with  $n$  denoting the total number of individuals. For a given prior distribution  $p(\boldsymbol{\theta})$ , the augmented posterior distribution of the model is given as

$$p(\boldsymbol{\theta}, \mathbf{z} | \mathbf{y}, X) \propto p(\mathbf{y}, \mathbf{z} | \boldsymbol{\theta}, X) p(\boldsymbol{\theta}). \quad (20)$$

Albert and Chib (1993) were the first to propose the use of data augmentation (see Section 3.1.3, herein) for conducting Bayesian analysis of the multinomial probit model.

**The predictive distribution.** Consider now an individual  $s$ , with predictor values  $X_s$ , whose choice behavior we would like to predict. The predictive for individual  $s$  can be written as

$$p(y_s | X_s, \mathbf{y}, X) = \int_{\Theta} \int_{\mathbf{z}_s} p(y_s | \mathbf{z}_s) p(\mathbf{z}_s | \boldsymbol{\theta}, X_s) d\mathbf{z}_s \int_{\mathbf{z}} p(\boldsymbol{\theta}, \mathbf{z} | \mathbf{y}, X) d\boldsymbol{\theta} d\mathbf{z}, \quad (21)$$

from which the predictive choice probabilities  $\Pr(Y_s = j | X_s, \mathbf{y}, X) = p(j | X_s, \mathbf{y}, X)$  can be constructed. The specification and computation of the predictive distribution in (21) poses three key challenges.

First,  $p(y_s | \mathbf{z}_s)$  requires a choice of base category. This choice affects the prior predictive choice probabilities, and hence the (posterior) predictive choice probabilities can be sensitive to the choice of base category; see Burgette and Nordheim (2012). Burgette, Puelz, and Hahn (2021) propose a symmetric prior specification to address this problem. The parameters  $\boldsymbol{\theta}$  are not identified under this prior, but this does not affect the predicted probabilities.

Second, the parameters  $\boldsymbol{\theta}$  lack scale identification, as  $p(y_i | \mathbf{z}_i) = p(y_i | c\mathbf{z}_i)$  for any positive scalar  $c$ . Different solutions have been proposed to fix the scale, all based on a constraint on the specification of  $\boldsymbol{\Sigma}$ . For instance, McCulloch, Polson, and Rossi (2000) fix the first leading element of  $\boldsymbol{\Sigma}$  to unity. This approach is sensitive to the ordering of the choice categories in the model. Burgette and Nordheim (2012) fix the trace of  $\boldsymbol{\Sigma}$ , which is invariant to the way in which the choice categories enter the model.

Third, the computation of  $p(y_s | X_s, \mathbf{y}, X)$  involves the evaluation of the integrals over the latent utilities in  $\mathbf{z}_s$  and  $\mathbf{z}$ . Since no analytical solution for these integrals is available, they are solved with MCMC sampling steps. The latent utility of each choice category is sampled from a univariate truncated normal, conditional on the latent utilities for all the other choice alternatives, for each individual (McCulloch & Rossi, 1994). Conditional on the draws for the latent utilities, sampling  $\boldsymbol{\beta}$  from its full conditional is straightforward. Generating from the conditional distribution of  $\boldsymbol{\Sigma}$  is non-standard, as the scale restrictions on  $\boldsymbol{\Sigma}$  have to be taken into account.

**Scalable Bayesian prediction.** In addition to the challenges delineated above, it is difficult to scale  $p(y_s | X_s, \mathbf{y}, X)$  to problems with large choice sets or a large number of observations. Recent advances in the computation of the predictive have focused on tackling the scalability issues in  $J$  and  $n$ , as we discuss below.

When considering a full covariance matrix specification for  $\boldsymbol{\Sigma}$ , the total number of parameters increases quadratically with  $J$ . For problems with large choice sets and small samples, this implies that the ratio of the total number of parameters to the total number of observations is large, making it difficult to construct accurate predictions. Loaiza-Maya and Nibbering (2022b) propose a spherical transformation of the covariance matrix of the latent utilities that imposes a parsimonious factor structure and a trace restriction. As a result, the total

number of parameters grows only linearly with  $J$ . The authors demonstrate that this parsimonious structure leads to improved predictive performance over full covariance matrix specifications.

Additionally, as noted above, the construction of the predictive entails an evaluation of the integral over the latent utilities  $\mathbf{z}$ . Although MCMC is able to solve this integral, it does so by generating the utility vector for each individual from a multivariate truncated normal, which is a computationally costly exercise (McCulloch & Rossi, 1994; Botev, 2017). This renders MCMC algorithms impractical for problems where a large  $n$  is considered.

VB can be employed to tackle problems with large  $n$ . Adapting the generic descriptions of VB in Section 3.2.3 and Appendix A.7, the application of VB in this setting considers the class of approximating densities  $\mathcal{Q}$  with elements  $q_\lambda(\theta, \mathbf{z}) \in \mathcal{Q}$ , indexed by a variational parameter vector  $\lambda$ . The exact augmented posterior is approximated by  $q_{\hat{\lambda}}(\theta, \mathbf{z})$  with an optimal variational parameter vector equal to

$$\hat{\lambda} = \underset{\lambda \in \Lambda}{\operatorname{argmin}} \operatorname{KL}[q_\lambda(\theta, \mathbf{z})|p(\theta, \mathbf{z}|\mathbf{y}, X)], \quad (22)$$

where  $\operatorname{KL}$  denotes the Kullback–Leibler divergence. The variational predictive is then constructed as

$$\hat{p}_\lambda(y_s|X_s, \mathbf{y}, X) = \int_{\Theta} \int_{\mathcal{Z}_s} p(y_s|\mathbf{z}_s)p(\mathbf{z}_s|\theta, X_s)d\mathbf{z}_s \int_{\mathcal{Z}} q_\lambda(\theta, \mathbf{z})d\mathbf{z}d\theta.$$

Calibrating the variational approximation requires a scale-identified expression for  $p(\mathbf{y}, \mathbf{z}|\theta, X)$ . To achieve this, Girolami and Rogers (2006) consider an identity matrix covariance structure, while Fasano and Durante (2022) fix  $\Sigma$  at predetermined values. Loaiza-Maya and Nibbeling (2022a) propose a method for a multinomial probit model with a factor covariance structure. This method uses the hybrid variational approximation  $q_\lambda(\theta, \mathbf{z}) = q_\lambda(\theta)p(\mathbf{z}|\mathbf{y}, \theta, X)$  introduced by Loaiza-Maya et al. (2022).

#### 4.4. Electricity pricing and demand (Anastasios Panagiotelis)

Forecasting in electricity markets is critical for the efficient day-to-day operation of power grids, for long-term planning of infrastructure, and increasingly, at a disaggregated level, for the management of smart grids. This section covers forecasting electricity prices, electricity load/demand, and generation by source of power, primarily wind and solar. Hereafter, these problems are collectively referred to as ‘electricity forecasting’. Motivations for electricity forecasting can be found in general reviews such as Weron (2014) for price forecasting, Lindberg, Seljom, Madsen, Fischer, and Korpås (2019) for load forecasting, Antonanzas et al. (2016) for solar power forecasting, and Giebel and Kariniotakis (2017) for wind power forecasting. These reviews indicate that the majority of work in electricity forecasting does not employ a Bayesian approach. Notwithstanding this, Bayesian methods have found success in the field.

There are very few instances of Bayesian forecasting in electricity markets that predate the early 2000s, although we here cover some notable exceptions. Bunn (1980) consider the case of updating load forecasts in an

online fashion by computing a Bayesian model average of load profiles of a cloudy and a sunny day. Meanwhile, Bayesian VARs have been used by Gunel (1987), Beck and Solow (1994), and Joutz, Maddala, and Trost (1995) to forecast energy demand, nuclear power generation, and demand prices and consumption, respectively. A Bayesian VAR shrinks autoregressive coefficients to either a random walk or white noise depending on whether data are stationary or non-stationary, and was popularized in macroeconomics by Doan et al. (1984) (see also Section 4.1). The performance of Bayesian VARs in early electricity forecasting applications is mixed: Beck and Solow (1994) find evidence in favor of Bayesian autoregression; Joutz et al. (1995) find that Bayesian VARs are effective for forecasting demand, but not price; while Gunel (1987) does not find any improvement at all from using Bayesian VARs rather than conventional autoregressive integrated moving average (ARIMA) models.

With the advent and popularization of MCMC methods, Bayesian forecasting has begun to find greater success in the field of electricity forecasting. In the literature of roughly the past two decades, there are three common major motivations for using Bayesian forecasting, two of which have antecedents in the earlier literature. The first is the use of Bayesian models,<sup>18</sup> which have now grown well beyond Bayesian VARs to include models with latent volatilities, models with a spatial dimension, and Bayesian neural networks. The second is the use of BMA for forecast combination. The third is the production of full probabilistic forecasts via Bayesian computation. These are now each discussed in turn.

*Bayesian models.* The structure inherent in many electricity forecasting problems provides a motivation for the innovative use of priors to improve forecasting accuracy. Although the early literature cited above found somewhat ambiguous results when comparing Bayesian VARs to classical alternatives, more recent work finds evidence in favor of a Bayesian approach; see Raviv, Bouwman, and Van Dijk (2015) for point forecasts and Gianfreda, Ravazzolo, and Rossini (2020) for both point and density forecasts. An important aspect of this work is the exploitation of the intraday nature of the data, since hourly prices are typically stacked in a VAR model. The intraday structure lends itself to priors that shrink parameters corresponding to consecutive hours of the day that are close to one another. An early application of this approach can be seen in Cottet and Smith (2003).

Since electricity data are increasingly available not only at a high temporal frequency but also at a high spatial resolution, there are further examples in the literature of using priors to exploit neighborhood structures. Examples include Ohtsuka, Oga, and Kakamu (2010), who use spatial ARMA processes to predict electricity load in nine Japanese regions, and Gilanifar, Wang, Ozguven, Zhou, and Arghandeh (2019), who use spatiotemporal

<sup>18</sup> By a ‘Bayesian model’ we generally mean a model with a prior and likelihood estimated by Bayesian inference. Bayesian methods for finding tuning parameters, such as the automatic relevance determination (see Hippert & Taylor, 2010, for an example in electricity forecasting), and Bayesian optimization lie beyond the scope of this section.

Gaussian processes to forecast residential-level electricity demand. Even where spatial information is unavailable, hierarchical models estimated using Bayesian methods have been used to produce disaggregate energy demand forecasts; examples can be found in [Mori and Nakano \(2014\)](#) and [Wang, Sun, and Lall \(2017\)](#), who use Gaussian processes, and in [Grillone et al. \(2021\)](#), who use regression. Informative hierarchical priors have been used in instances where data sets are small in size or unavailable; for example, [Pezzulli et al. \(2006\)](#) elicit priors for future trajectories of temperature in the winter using past observations, and [Launay, Philippe, and Lamarche \(2015\)](#) elicit priors for the electricity demand of non-metered households using data on metered households.

While the aforementioned examples take a Bayesian approach to exploit the use of priors in novel ways, another strain of the Bayesian forecasting literature is based on estimating models with latent variables. Examples in electricity forecasting include a latent jump process for price spikes ([Chan, Choy, & Lam, 2014](#)) and SV models ([Smith, 2010](#); [Kostrzewski & Kostrzewska, 2019](#)). Also, in recent years, Bayesian analysis of machine learning models has become increasingly popular. This includes neural network models ([Brusaferri, Matteucci, Portolani, & Vitali, 2019](#); [Hayekhloo, Azimi, Ghofrani, Menhaj, & Shekari, 2019](#); [Capone, Helming, & Hirche, 2020](#)), where VB is typically used. Also, Bayesian regression trees (see Section 4.1) have been applied to electricity forecasting by [Nateghi, Guikema, and Quiring \(2011\)](#) and [Alipour, Mukherjee, and Nateghi \(2019\)](#), who find that they outperform non-Bayesian counterparts. Finally, there is extensive literature on using Bayesian networks for forecasting in energy; see [Adedipe, Shafiee, and Zio \(2020\)](#) for a review of these methods in forecasting wind generation.

*Bayesian model averaging (BMA).* As noted above, the importance of forecast combination is widely appreciated in the forecasting literature. As highlighted in Section 3.3.3, many different Bayesian approaches to forecast combination have now been explored, yet BMA remains a very important method in the sphere of electricity forecasting. As described in Section 2.1, BMA uses posterior model probabilities as combination weights. Whenever the choice of model is parameterized, the predictive density has an interpretation as a forecast combination. Examples include [Smith \(2000\)](#), who combines forecasts from regression models that include different predictor sets, and [Panagiotelis and Smith \(2008\)](#), who average over models with different combinations of skew and symmetric marginal distributions.

It is also common in the electricity forecasting literature to produce point forecasts from different models and then combine these using BMA as a post-processing step. This approach grew out of research combining ensembles of forecasts from numerical weather predictions (NWP) ([Raftery, Gneiting, Balabdaoui, & Polakowski, 2005](#); [Sloughter, Gneiting, & Raftery, 2010](#)). In the NWP setting, forecasts are the outputs of deterministic physical models. Statistical models are then formed by assuming that for  $k = 1, \dots, K$ ,  $p(y_t | a_k, b_k, f_k, \sigma^2, \mathcal{M}_k) \sim N(a_k + b_k f_k, \sigma^2)$ , where  $f_k$  is the  $k$ th NWP, and  $a_k, b_k$ , and  $\sigma^2$  are

additional parameters. These statistical models are then combined using the usual BMA machinery described by (3), with the key distinction being that posterior model probabilities are replaced with  $p(M_k | y_{T-L+1:T})$ , where  $L$  is the length of the window. Uncertainty over  $a_k, b_k$  and  $\sigma^2$  is integrated out in the usual way, and there are no additional parameters, since the  $f_k$  are obtained deterministically. This approach has been used in energy forecasting by [Antonanzas et al. \(2006\)](#), who motivate forecasting rainfall as a input into forecasting generation from hydroelectric dams, and [Du \(2018\)](#), who uses wind forecasts to predict generation from wind farms.

The work of [Raftery et al. \(2005\)](#) has been subsequently extended to the case where the forecasts  $f_k$  are not the outputs of deterministic physical models, but are rather point forecasts from statistical models, each with its own unknown parameters. For example [Nowotarski, Raviv, Trück, and Weron \(2014\)](#) adopt the approach of [Raftery et al. \(2005\)](#) but where the  $f_k$  are obtained from statistical time series models with parameters estimated using frequentist techniques. This approach is not fully Bayesian (despite being referred to as BMA in the literature), since although the model average integrates over the uncertainty in  $a_k, b_k$  and  $\sigma^2$ , it does not integrate over uncertainty in the parameters of the underlying time series models used to generate the point forecasts  $f_k$ .<sup>19</sup> In a similar vein, [Hassan, Khosravi, and Jaafar \(2015\)](#) and [Raza, Nadarajah, and Ekanayake \(2017\)](#) combine electricity load forecasts from different neural networks.

*Probabilistic forecasting.* A common motivation for taking a Bayesian approach is the ease with which the computational machinery of MCMC or approximate methods produces a full predictive density rather than only point forecasts. Key operational decisions in electricity forecasting depend on quantities other than the predicted mean; see [Nowotarski and Weron \(2018\)](#) and the references therein for discussion. While the importance of probabilistic forecasting is often highlighted in Bayesian papers, it is not always the case that forecasts are evaluated in a way that assesses the quality of the full predictive distribution.<sup>20</sup> For example, probabilistic forecasts are often summarized by prediction intervals, and the empirical coverage of these intervals is used as a means of checking model quality; for an early example, see [Pezzulli et al. \(2006\)](#), and more recently [Wang et al. \(2017\)](#) and [Kostrzewski and Kostrzewska \(2019\)](#), where the latter show that Bayesian methods compare favorably to non-Bayesian alternatives for forecasting electricity prices. [Kostrzewski and Kostrzewska \(2019\)](#) also evaluate  $\alpha$ -level quantile forecasts  $\hat{q}_t$  using the pinball loss,

$$L_\alpha(y_t, \hat{q}_t) = \alpha(y_t - \hat{q}_t)I[y_t \geq \hat{q}_t] + (1 - \alpha)(\hat{q}_t - y_t)I[y_t < \hat{q}_t].$$

<sup>19</sup> The same point does not apply when combining ensembles from NWPs, since the forecasting models are deterministic.

<sup>20</sup> We note that in some cases this is challenging, for example with long-run forecasts, as in [Da Silva, Oliveira, and Souza \(2019\)](#).

Yang, Li, Gulliver, and Li (2019) and Sun, Zhang, Wang, Strbac, and Kang (2019) also use the pinball loss to evaluate forecasts of residential-level load (net of solar PV generation in the latter case).

However, the use of scoring rules (Gneiting & Raftery, 2007), and hence the explicit recognition of the distributional form of the forecasts, is becoming increasingly popular as a means of evaluating predictive distributions in both Bayesian and non-Bayesian electricity forecasting. The continuously ranked probability score (Gneiting & Raftery, 2007) is particularly amenable to Bayesian inference, since it is usually approximated using a Monte Carlo sample from the predictive density. For an early example of its use in Bayesian electricity forecasting, see Panagiotelis and Smith (2008); for later examples, see Bracale and De Falco (2015), Brusaferrri et al. (2019), and Gianfreda et al. (2020). Other scoring rules are less commonly used in the Bayesian electricity forecasting literature, although Ohtsuka et al. (2010), where the log score is used, is a notable exception.

## 5. In summary

Bayesian forecasting is underpinned by a single core principle: Uncertainty about the future value of a random variable is expressed using a probability distribution, where the form of that distribution in turn reflects uncertainty about all other unknowns on which the investigator chooses not to condition.

While this principled approach to forecasting is arguably one of the most compelling features of the paradigm, the challenge has potentially been in the *implementation* of Bayesian forecasting: namely, computing the expectation that defines the predictive distribution, particularly when it is difficult to access (draws from) the posterior itself. And as models have become larger and more challenging, and as data sets have grown bigger, this problem of accessing the exact posterior has only increased. However, as this review has demonstrated, the expansion of the forecasting problems being tackled has gone hand-in-hand with the development of new and improved computational methods designed expressly to access challenging posteriors, and in a reasonable computing time. Notably, when it comes to accurate forecasting, somewhat crude approximations of the posterior have been found to still yield accurate predictions, meaning that Bayesian forecasting remains viable for large and complex models for which approximate computation of posteriors is the only feasible approach.

The more fundamental problem of model misspecification can also be managed by moving away from the conventional likelihood-based Bayesian updating and allowing forecast accuracy itself – and its link to the future decisions that depend on that accuracy – to drive the updating. This, in turn, ensures that forecasts are ‘fit for purpose’, *despite* the inevitable misspecification of the forecasting model. Allied with the computational power that now drives the Bayesian engine, this ability to generalize the paradigm beyond its traditional links with the likelihood principle is a potent, if not yet fully realized, force in forecasting.

## Declaration of competing interest

The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

## Appendix. Further computational details

### A.1. Gibbs sampling

Under the required regularity conditions (see Tierney, 1994), the Gibbs sampler yields a Markov chain with invariant distribution,  $p(\theta|\mathbf{y}_{1:T})$ , via a transition kernel that is defined as the product of full conditional posteriors associated with the joint. For the case of  $\theta$  partitioned into  $B$  mutually exclusive blocks,  $\theta = (\theta'_1, \theta'_2, \dots, \theta'_b, \dots, \theta'_B)'$ , the steps of the Gibbs algorithm are given in Algorithm 1.

### A.2. MH-within-Gibbs sampling

In Algorithm 2 we provide the generic steps of the so-called MH-within-Gibbs algorithm for the case of  $\theta$  partitioned into  $B$  mutually exclusive blocks,  $\theta = (\theta'_1, \theta'_2, \dots, \theta'_b, \dots, \theta'_B)'$ . The symbol  $p_b^*$  represents (the ordinate of) a kernel of the corresponding conditional  $p_b(\cdot|\cdot)$ .

The  $b$ th candidate density  $q_b(\theta_b|\mathbf{y}_{1:T})$  may be chosen to deliberately target the form of the  $b$ th conditional density,  $p_b(\theta_b|\theta^{(i)}_1, \dots, \theta^{(i)}_{b-1}, \theta^{(i-1)}_{b+1}, \dots, \theta^{(i-1)}_B, \mathbf{y}_{1:T})$ , in which case the algorithm may be referred to as a ‘tailored’ algorithm; otherwise  $q_b(\theta_b|\mathbf{y}_{1:T})$  may be chosen in a more automated fashion, such as in a random-walk MH algorithm. The references cited in the text provide all details.

### A.3. MH-within-Gibbs sampling in state-space models

The application of MH-within-Gibbs sampling within a state-space setting is qualitatively the same as described in Algorithm 2, except that the joint set of unknowns is augmented to  $(\theta, \mathbf{z}_{1:T})$ , and decisions about partitioning need to be made for both  $\theta|\mathbf{z}_{1:T}$  and  $\mathbf{z}_{1:T}|\theta$ . Decisions about the blocking of  $\mathbf{z}_{1:T}$  are particularly important, given both the dimension of  $\mathbf{z}_{1:T}$  and the time-series dependence in the state process, as are matters of parameterizing the state-space model. We refer the reader to Shephard and Pitt (1997) and Strickland et al. (2006) for illustrations of state blocking in which the block sizes are selected randomly, and to Frühwirth-Schnatter (2004) and Strickland et al. (2008) for explorations of the impact of parameterization on the performance of the sampler.

### A.4. PMMH in state-space models

Early Bayesian treatments of non-linear state-space models often exploited a linear Gaussian approximation at some point, for the purpose of defining candidate densities for (blocks of)  $\mathbf{z}_{1:T}$  (e.g. Kim et al., 1998; Stroud et al., 2003; Strickland et al., 2006), thereby enabling a Kalman filter-based forward filtering, backward sampling algorithm (Carter & Kohn, 1994; Frühwirth-Schnatter, 1994)



**Algorithm 1** Gibbs sampling algorithm.

---

Specify an initial value  $\theta^{(0)}$  and partition the parameter set into  $B$  mutually exclusive blocks  
**for**  $i = 1, \dots, M$  **do**  
  **for**  $b = 1, \dots, B$  **do**  
    Draw  $\theta_b^{(i)} \sim p_b(\theta_b | \theta_1^{(i)}, \dots, \theta_{b-1}^{(i)}, \theta_{b+1}^{(i-1)}, \dots, \theta_B^{(i-1)}, \mathbf{y}_{1:T})$   
  **end for**  
**end for**  
Return a sample of draws from  $p(\theta | \mathbf{y}_{1:T})$ .

---

**Algorithm 2** MH-within-Gibbs algorithm.

---

Specify an initial value  $\theta^{(0)}$ , a partition of the parameter set into  $B$  mutually exclusive blocks, and a proposal distribution  $q_b(\theta_b | \mathbf{y}_{1:T})$  for  $b \in \{1, \dots, B\}$ .  
**for**  $i = 1, \dots, M$  **do**  
  **for**  $b = 1, \dots, B$  **do**  
    Draw  $\theta_b^c \sim q_b(\theta_b | \mathbf{y}_{1:T})$   
    Compute the Metropolis–Hastings ratio:  

$$r = \frac{p_b^*(\theta_b^c | \theta_1^{(i)}, \dots, \theta_{b-1}^{(i)}, \theta_{b+1}^{(i-1)}, \dots, \theta_B^{(i-1)}, \mathbf{y}_{1:T}) \times q_b(\theta_b^{(i-1)} | \mathbf{y}_{1:T})}{p_b^*(\theta_b^{(i-1)} | \theta_1^{(i)}, \dots, \theta_{b-1}^{(i)}, \theta_{b+1}^{(i-1)}, \dots, \theta_B^{(i-1)}, \mathbf{y}_{1:T}) \times q_b(\theta_b^c | \mathbf{y}_{1:T})}$$
  
    **if**  $\mathcal{U}(0, 1) < r$  **then**  
      Set  $\theta_b^{(i)} = \theta_b^c$   
    **else**  
      Set  $\theta_b^{(i)} = \theta_b^{(i-1)}$   
    **end if**  
  **end for**  
**end for**  
Return a sample of draws from  $p(\theta | \mathbf{y}_{1:T})$ .

---

to be used to produce a candidate draw of (any particular block of)  $\mathbf{z}_{1:T}$ , conditional on  $\theta$ . As noted in Section 3.2.2 (and in the review by Giordani et al., 2011), more recent approaches to such models have exploited PMMH principles instead. Algorithm 3 reproduces the algorithm in Andrieu et al. (2011) (Section 2.4.2 therein), adapted slightly to match the notation of the current paper. To simplify the exposition, the algorithm is presented for sampling the full vector  $\theta$ . In practice the algorithm would be modified to cater for any blocking of  $\theta$ .

**A.5. ABC based on summary statistics**

The simplest (accept/reject) form of the ABC algorithm, as based on a chosen vector of summaries,  $\eta(\mathbf{y}_{1:T})$ , proceeds via the steps in Algorithm 4, with the accepted draws of  $\theta$  used to produce an estimate of  $p_\varepsilon(\theta | \eta(\mathbf{y}_{1:T}))$  via kernel density methods. This posterior is equivalent to  $p(\theta | \mathbf{y}_{1:T})$  if and only if  $\eta(\mathbf{y}_{1:T})$  is sufficient for conducting inference on  $\theta$ , and for  $\varepsilon \rightarrow 0$ . Clearly, the very problems for which ABC is required imply that sufficient statistics are not available, and the requirement that  $\varepsilon \rightarrow 0$  is infeasible in practice; so inference via ABC is only ever intrinsically approximate.<sup>21</sup>

<sup>21</sup> The notation  $\mathbf{z}_{1:T}$  used in this section and in Appendix A.6, below, is not to be confused with the use of  $\mathbf{z}_{1:T}$  to denote a vector of latent variables elsewhere in the paper.

**A.6. BSL based on summary statistics**

BSL mimics ABC in targeting a posterior for  $\theta$  that conditions on a vector of summaries  $\eta(\mathbf{y}_{1:T})$ , rather than the full data set  $\mathbf{y}_{1:T}$ . However, the summaries play a different role in the algorithm. Once again with reference to the simplest version of the algorithm, the steps of the BSL-MCMC algorithm are as given in Algorithm 5. Note that for a given  $\theta$ , the draws  $\mathbf{z}_{(j)1:T} \sim i.i.d. p(\cdot | \theta)$ ,  $j = 1, \dots, m$ , are used to estimate  $\mu(\theta)$  and  $\Sigma(\theta)$  as  $\mu_m(\theta) = \frac{1}{m} \sum_{j=1}^m \eta(\mathbf{z}_{(j)1:T})$  and  $\Sigma_m(\theta) = \frac{1}{m-1} \sum_{j=1}^m (\eta(\mathbf{z}_{(j)1:T}) - \mu_m(\theta))(\eta(\mathbf{z}_{(j)1:T}) - \mu_m(\theta))'$ . The  $M$  draws of  $\theta$  are used to produce an estimate of  $p(\theta | \eta(\mathbf{y}_{1:T}))$  via kernel density methods.

**A.7. VB**

VB seeks the best approximation to  $p(\theta | \mathbf{y}_{1:T})$  over a ‘variational family’ of densities  $\mathcal{Q}$ , with generic element  $q(\theta)$ . Typically, this proceeds by minimizing the Kullback–Leibler (KL) divergence between  $q(\theta)$  and  $p(\theta | \mathbf{y}_{1:T})$ , which produces the variational approximation as

$$q^*(\theta) := \operatorname{argmin}_{q(\theta) \in \mathcal{Q}} \text{KL}[q(\theta) | p(\theta | \mathbf{y}_{1:T})], \quad (23)$$

where

$$\text{KL}[q(\theta) | p(\theta | \mathbf{y}_{1:T})] = \mathbb{E}_q[\log(q(\theta))] - \mathbb{E}_q[\log(p(\theta, \mathbf{y}_{1:T}))] + \log(p(\mathbf{y}_{1:T})) \quad (24)$$

**Algorithm 3** PMMH algorithm.

Step 1: Initialization,  $i = 0$

- (a) Set  $\theta^{(0)}$  arbitrarily and
- (b) Run an SMC algorithm targeting  $p(\mathbf{z}_{1:T}|\mathbf{y}_{1:T}, \theta^{(0)})$ , sample  $\mathbf{z}_{1:T}^{(0)} \sim \hat{p}(\mathbf{z}_{1:T}|\mathbf{y}_{1:T}, \theta^{(0)})$ , and let  $\hat{p}(\mathbf{y}_{1:T}|\theta^{(0)})$  denote the marginal likelihood estimate.

Step 2:

**for**  $i = 1, \dots, M$  **do**

- (a) Draw  $\theta^c \sim q(\theta|\mathbf{y}_{1:T}, \theta^{(i-1)})$ ,

- (b) Run an SMC algorithm targeting  $p(\mathbf{z}_{1:T}|\mathbf{y}_{1:T}, \theta^c)$ , sample  $\mathbf{z}_{1:T}^c \sim \hat{p}(\mathbf{z}_{1:T}|\mathbf{y}_{1:T}, \theta^c)$ , and let  $\hat{p}(\mathbf{y}_{1:T}|\theta^c)$  denote the marginal likelihood estimate.

- (c) Compute the Metropolis–Hastings ratio:

$$r = \frac{\hat{p}(\mathbf{y}_{1:T}|\theta^c)p(\theta^c) \times q(\theta^{(i-1)}|\mathbf{y}_{1:T}, \theta^c)}{\hat{p}(\mathbf{y}_{1:T}|\theta^{(i-1)})p(\theta^{(i-1)}) \times q(\theta^c|\mathbf{y}_{1:T}, \theta^{(i-1)})}$$

**if**  $\mathcal{U}(0, 1) < r$  **then**

Set  $\theta^{(i)} = \theta^c$ ,  $\mathbf{z}_{1:T}^{(i)} = \mathbf{z}_{1:T}^c$ ,  $\hat{p}(\mathbf{y}_{1:T}|\theta^{(i)}) = \hat{p}(\mathbf{y}_{1:T}|\theta^c)$

**else**

Set  $\theta^{(i)} = \theta^{(i-1)}$ ,  $\mathbf{z}_{1:T}^{(i)} = \mathbf{z}_{1:T}^{(i-1)}$ ,  $\hat{p}(\mathbf{y}_{1:T}|\theta^{(i)}) = \hat{p}(\mathbf{y}_{1:T}|\theta^{(i-1)})$

**end if**

**end for**

Return a sample of draws from  $p(\theta|\mathbf{y})$ .

**Algorithm 4** Accept/reject ABC algorithm based on summary statistics.

**for**  $i = 1, \dots, M$  **do**

Simulate  $\theta^{(i)}$ ,  $i = 1, 2, \dots, M$ , from  $p(\theta)$ , and artificial data  $\mathbf{z}_{1:T}^{(i)}$  from  $p(\cdot|\theta^{(i)})$ ;

Accept  $\theta^{(i)}$  if  $d\{\eta(\mathbf{z}_{1:T}^{(i)}), \eta(\mathbf{y}_{1:T})\} \leq \varepsilon$ , where  $d\{\cdot, \cdot\}$  denotes a generic metric and  $\varepsilon > 0$  a pre-specified tolerance parameter.

**end for**

and  $p(\theta, \mathbf{y}_{1:T}) = p(\mathbf{y}_{1:T}|\theta)p(\theta)$ . Given that the unknown normalizing constant  $\log(p(\mathbf{y}_{1:T}))$  in (24) does not depend on  $q$ , the (infeasible) optimization problem in (23) is replaced by the following equivalent (and feasible) optimization problem:

$$q^*(\theta) := \operatorname{argmax}_{q(\theta) \in \mathcal{Q}} \left\{ \mathbb{E}_q[\log(p(\theta, \mathbf{y}_{1:T}))] - \mathbb{E}_q[\log(q(\theta))] \right\}, \quad (25)$$

with the so-called evidence lower bound (ELBO) defined as:

$$\text{ELBO}[q(\theta)] := \mathbb{E}_q[\log(p(\theta, \mathbf{y}_{1:T}))] - \mathbb{E}_q[\log(q(\theta))]. \quad (26)$$

The usefulness of VB is that, for certain models,  $p(\mathbf{y}_{1:T}|\theta)$ , and certain choices of  $\mathcal{Q}$ , the optimization problem in (25) can be solved efficiently using various numerical

**Algorithm 5** BSL-MCMC algorithm.

**for**  $i = 1, \dots, M$  **do**

Draw  $\theta^* \sim q(\theta|\theta^{(i-1)})$

Produce  $\mu_m(\theta)$  and  $\Sigma_m(\theta)$  using  $j = 1, \dots, m$  independent model simulations at  $\theta^*$

Compute the synthetic likelihood  $L^* = \mathcal{N}[\eta(\mathbf{y}); \mu_m(\theta^*), \Sigma_m(\theta^*)]$  and  $L^{(i-1)}$  defined in a corresponding manner

Compute the Metropolis–Hastings ratio:

$$r = \frac{L^* \pi(\theta^*) q(\theta^{(i-1)}|\theta^*)}{L^{(i-1)} \pi(\theta^{(i-1)}) q(\theta^*|\theta^{(i-1)})}$$

**if**  $\mathcal{U}(0, 1) < r$  **then**

Set  $\theta^{(i)} = \theta^*$ ,  $\mu_m(\theta^{(i)}) = \mu_m(\theta^*)$ , and  $\Sigma_m(\theta^{(i)}) = \Sigma_m(\theta^*)$

**else**

Set  $\theta^{(i)} = \theta^{(i-1)}$ ,  $\mu_m(\theta^{(i)}) = \mu_m(\theta^{(i-1)})$ , and  $\Sigma_m(\theta^{(i)}) = \Sigma_m(\theta^{(i-1)})$ .

**end if**

**end for**

algorithms. Most notably, for problems in which the dimension of the unknowns, and possibly that of  $\mathbf{y}_{1:T}$ , is large, the production of  $q^*(\theta)$  is much faster (often orders of magnitude so) than producing an estimate of  $p(\theta|\mathbf{y}_{1:T})$  (and any associated quantities, including predictives) via simulation (Braun & McAuliffe, 2010; Kabisa, Dunson, & Morris, 2016; Wand, 2017; Koop & Korobilis, 2023). The relationship between (24) and (26), plus the fact that  $KL[\cdot] \geq 0$ , also means  $\text{ELBO}[q^*(\theta)]$  is a lower bound on the logarithm of the ‘evidence’, or marginal likelihood,  $p(\mathbf{y}_{1:T})$ , hence the abbreviation ‘ELBO’.

Different VB methods are defined by both the choice of  $\mathcal{Q}$  and the manner in which the optimization is implemented, and we refer the reader to Ormerod and Wand (2010), Blei et al. (2017), and Zhang, Bütetage, Kjellström, and Mandt (2018) for reviews, including algorithmic details for specific VB methods.

**A.8. INLA**

Rue et al. (2009) adapted the very early approximation method of Laplace (1774) to approximate posteriors (and associated quantities) in the latent Gaussian model class, which encompasses a large range of (potentially high-dimensional) models, including the non-Gaussian state-space models that feature heavily in economics and finance. In brief, Rue et al. use a series of nested Laplace approximations, allied with low-dimensional numerical integration, thus calling their method integrated nested Laplace approximation (INLA). As with VB, INLA eschews simulation for optimization, exploiting bespoke numerical algorithms designed for the specific (albeit broad) model class. We refer the reader to Rue et al. (2009), Rue et al. (2017), Martino and Riebler (2019), and Wood (2019) for implementation details.

## References

- Aastveit, K. A., Cross, J. L., & van Dijk, H. K. (2023). Quantifying time-varying forecast uncertainty and risk for the real price of oil. *Journal of Business & Economic Statistics*, 41(2), 523–527.
- Aastveit, K. A., Ravazzolo, F., & van Dijk, H. K. (2018). Combined density nowcasting in an uncertain economic environment. *Journal of Business & Economic Statistics*, 36(1), 131–145.
- Adedipe, T., Shafiee, M., & Zio, E. (2020). Bayesian network modelling for the wind energy industry: An overview. *Reliability Engineering & System Safety*, 202, Article 107053.
- Adolfson, M., Lindé, J., & Villani, M. (2007). Forecasting performance of an open economy DSGE model. *Econometric Reviews*, 26(2–4), 289–328.
- Adrian, T., Boyarchenko, N., & Giannone, D. (2021). Multimodality in macrofinancial dynamics. *International Economic Review*, 62(2), 861–886.
- Albert, J. H., & Chib, S. (1993). Bayesian analysis of binary and polychotomous response data. *Journal of the American Statistical Association*, 88(422), 669–679.
- Alexopoulos, A., Dellaportas, P., & Papaspiliopoulos, O. (2022). Bayesian prediction of jumps in large panels of time series data. *Bayesian Analysis*, 17(2), 651–683.
- Alipour, P., Mukherjee, S., & Nateghi, R. (2019). Assessing climate sensitivity of peak electricity load for resilient power systems planning and operation: A study applied to the Texas region. *Energy*, 185, 1143–1153.
- Andrieu, C., Doucet, A., & Holenstein, R. (2011). Particle Markov chain Monte Carlo. *The Journal of the Royal Statistical Society, Series B*, 72(2), 269–342, With discussion.
- Andrieu, C., Doucet, A., & Robert, C. (2004). Computational advances for and from Bayesian analysis. *Statistical Science*, 19(1), 118–127.
- Andrieu, C., & Roberts, G. (2009). The pseudo-marginal approach for efficient Monte Carlo computations. *The Annals of Statistics*, 37(2), 697–725.
- Ansari, A., Li, Y., & Zhang, J. Z. (2018). Probabilistic topic model for hybrid recommender systems: A stochastic variational Bayesian approach. *Marketing Science*, 37(6), 987–1008.
- Antonanzas, J., Osorio, N., Escobar, R., Urraca, R., Martinez-de Pison, F. J., & Antonanzas-Torres, F. (2016). Review of photovoltaic power forecasting. *Solar Energy*, 136, 78–111.
- Araya, S., Elberg, A., Noton, C., & Schwartz, D. (2022). Identifying food labeling effects on consumer behavior. *Marketing Science*, 41(5), 871–1027.
- Ardia, D., Baştürk, N., Hoogerheide, L., & van Dijk, H. K. (2012). A comparative study of Monte Carlo methods for efficient evaluation of marginal likelihood. *Computational Statistics & Data Analysis*, 56(11), 3398–3414.
- Ausín, M. C., Galeano, P., & Ghosh, P. (2014). A semiparametric Bayesian approach to the analysis of financial time series with applications to value at risk estimation. *European Journal of Operational Research*, 232(2), 350–358.
- Baştürk, N., Borowska, A., Grassi, S., Hoogerheide, L., & van Dijk, H. (2019). Forecast density combinations of dynamic models and data driven portfolio strategies. *Journal of Econometrics*, 210(1), 170–186.
- Baker, J., Fearnhead, P., Fox, E., & Nemeth, C. (2019). Control variates for stochastic gradient MCMC. *Statistics and Computing*, 29, 599–615.
- Bakshi, G., Cao, C., & Chen, Z. (1997). Empirical performance of alternative option pricing models. *The Journal of Finance*, 52(5), 2003–2049.
- Barbura, M., Giannone, D., & Reichlin, L. (2010). Large Bayesian vector autoregressions. *Journal of Applied Econometrics*, 25(1), 71–92.
- Bardenet, R., Doucet, A., & Holmes, C. (2017). On Markov chain Monte Carlo methods for tall data. *Journal of Machine Learning Research*, 18(1), 1515–1557.
- Basetti, F., Casarin, R., & Ravazzolo, F. (2018). Bayesian nonparametric calibration and combination of predictive distributions. *Journal of the American Statistical Association*, 113(522), 675–685.
- Bastani, H., Simchi-Levi, D., & Zhu, R. (2022). Meta dynamic pricing: Transfer learning across experiments. *Management Science*, 68(3), 1865–1881.
- Baştürk, N., Borowska, A., Grassi, S., Hoogerheide, L., & van Dijk, H. K. (2019). Forecast density combinations of dynamic models and data driven portfolio strategies. *Journal of Econometrics*, 210(1), 170–186.
- Bauwens, L., & Lubrano, M. (1998). Bayesian inference on GARCH models using the Gibbs sampler. *The Econometrics Journal*, 1(1), 23–46.
- Beaumont, M. (2003). Estimation of population growth or decline in genetically monitored populations. *Genetics*, 164(3), 1139–1160.
- Beck, R., & Solow, J. L. (1994). Forecasting nuclear power supply with Bayesian autoregression. *Energy Economics*, 16(3), 185–192.
- Bernanke, B. S., Boivin, J., & Elias, P. (2005). Measuring the effects of monetary policy: A factor-augmented vector autoregressive (FAVAR) approach. *Quarterly Journal of Economics*, 120(1), 387–422.
- Bernardi, M., Gayraud, G., & Petrella, L. (2015). Bayesian tail risk interdependence using quantile regression. *Bayesian Analysis*, 10(3), 553–603.
- Bernardo, J., & Smith, A. (1994). *Bayesian theory*. New York: John Wiley.
- Betancourt, M. (2018). A conceptual introduction to Hamiltonian Monte Carlo. arXiv preprint <https://arxiv.org/abs/1701.02434v2>.
- Billio, M., Casarin, R., Ravazzolo, F., & van Dijk, H. (2013). Time-varying combinations of predictive densities using nonlinear filtering. *Journal of Econometrics*, 177(2), 213–232.
- Bissiri, P. G., Holmes, C. C., & Walker, S. G. (2016). A general framework for updating belief distributions. *Journal of the Royal Statistical Society Series B*, 78(5), 1103–1130.
- Black, F., & Scholes, M. (1973). The pricing of options and corporate liabilities. *Journal of Political Economy*, 81(3), 637–659.
- Blei, D. M., Kucukelbir, A., & McAuliffe, J. D. (2017). Variational inference: A review for statisticians. *Journal of the American Statistical Association*, 112(518), 859–877.
- Bollerslev, T., Chou, R., & Kroner, K. (1992). ARCH modeling in finance. A review of the theory and empirical evidence. *Journal of Econometrics*, 52(1–2), 5–59.
- Botev, Z. I. (2017). The normal law under linear restrictions: Simulation and estimation via minimax tilting. *Journal of the Royal Statistical Society. Series B. Statistical Methodology*, 79(1), 125–148.
- Bracale, A., & De Falco, P. (2015). An advanced Bayesian method for short-term probabilistic forecasting of the generation of wind power. *Energies*, 8(9), 10293–10314.
- Braun, M., & McAuliffe, J. (2010). Variational inference for large-scale models of discrete choice. *Journal of the American Statistical Association*, 105(489), 324–335.
- Brusaferri, A., Matteucci, M., Portolani, P., & Vitali, A. (2019). A Bayesian deep learning-based method for probabilistic forecast of day-ahead electricity prices. *Applied Energy*, 250, 1158–1175.
- Bunn, D. W. (1980). Experimental study of a Bayesian method for daily electricity load forecasting. *Applied Mathematical Modelling*, 4(2), 113–116.
- Burgette, L. F., & Nordheim, E. V. (2012). The trace restriction: An alternative identification strategy for the Bayesian multinomial probit model. *Journal of Business & Economic Statistics*, 30(3), 404–410.
- Burgette, L. F., Puelz, D., & Hahn, P. R. (2021). A symmetric prior for multinomial probit models. *Bayesian Analysis*, 16(3), 991–1008.
- Calvet, L. E., & Czellar, V. (2015). Accurate methods for approximate Bayesian computation filtering. *Journal of Financial Econometrics*, 13(4), 798–838.
- Canale, A., & Ruggiero, M. (2016). Bayesian nonparametric forecasting of monotonic functional time series. *Electronic Journal of Statistics*, 10(2), 3265–3286.
- Capone, A., Helminger, C., & Hirche, S. (2020). Day-ahead scheduling of thermal storage systems using Bayesian neural networks. *IFAC-PapersOnLine*, 53(2), 13281–13286.
- Carriero, A., Clark, T. E., & Marcellino, M. (2015). Bayesian VARs: Specification choices and forecast accuracy. *Journal of Applied Econometrics*, 30(1), 46–73.
- Carriero, A., Clark, T. E., & Marcellino, M. (2016). Common drifting volatility in large Bayesian VARs. *Journal of Business & Economic Statistics*, 34(3), 375–390.
- Carriero, A., Clark, T. E., & Marcellino, M. (2019). Large Bayesian vector autoregressions with stochastic volatility and non-conjugate priors. *Journal of Econometrics*, 212(1), 137–154.
- Carriero, A., Kapetanios, G., & Marcellino, M. (2009). Forecasting exchange rates with a large Bayesian VAR. *International Journal of Forecasting*, 25(2), 400–417.
- Carter, C. K., & Kohn, R. (1994). On Gibbs sampling for state space models. *Biometrika*, 81(3), 541–553.

- Carverhill, A., & Luo, D. (2023). A Bayesian analysis of time-varying jump risk in S&P 500 returns and options. *Journal of Financial Markets*, 64, Article 100786.
- Casarin, R., Grassi, S., Ravazzolo, F., & van Dijk, H. (2015). Parallel sequential Monte Carlo for efficient density combination: The DeCo MATLAB toolbox. *Journal of Statistical Software, Articles*, 68(3), 1–30.
- Casarin, R., Grassi, S., Ravazzolo, F., & van Dijk, H. K. (2023). A flexible predictive density combination for large financial data sets in regular and crisis periods. *Journal of Econometrics*, <https://doi.org/10.1016/j.jeconom.2022.11.004>.
- Casarin, R., Leisen, F., Molina, G., & ter Horst, E. (2015). A Bayesian beta Markov random field calibration of the term structure of implied risk neutral densities. *Bayesian Analysis*, 10(4), 791–819.
- Casarin, R., Mantoan, G., & Ravazzolo, F. (2016). Bayesian calibration of generalized pools of predictive distributions. *Econometrics*, 4(1), 1–24.
- Casella, G., & George, E. (1992). An introduction to Gibbs sampling. *American Statistician*, 46(3), 167–174.
- Ceruzzi, P. (2003). *A history of modern computing* (2nd ed.). MIT Press.
- Chakraborty, A., Nott, D. J., Drovandi, C., Frazier, D. T., & Sisson, S. A. (2022). Modularized Bayesian analyses and cutting feedback in likelihood-free inference. arXiv preprint arXiv:2203.09782.
- Chan, J. C. (2021). Minnesota-type adaptive hierarchical priors for large Bayesian VARs. *International Journal of Forecasting*, 37(3), 1212–1226.
- Chan, J. (2022). Asymmetric conjugate priors for large Bayesian VARs. *Quantitative Economics*, 13(3), 1145–1169.
- Chan, J. S., Choy, S. B., & Lam, C. P. (2014). Modeling electricity price using a threshold conditional autoregressive geometric process jump model. *Communications in Statistics. Theory and Methods*, 43(10–12), 2505–2515.
- Chan, J. C., Eisenstat, E., & Strachan, R. W. (2020). Reducing the state space dimension in a large TVP-VAR. *Journal of Econometrics*, 218(1), 105–118.
- Chan, J. C., Koop, G., & Potter, S. M. (2013). A new model of trend inflation. *Journal of Business & Economic Statistics*, 31(1), 94–106.
- Chan, J. C., Koop, G., & Yu, X. (2023). Large order-invariant Bayesian VARs with stochastic volatility. *Journal of Business and Economic Statistics*, (Forthcoming).
- Chan, J. C., & Yu, X. (2022). Fast and accurate variational inference for large Bayesian VARs with stochastic volatility. *Journal of Economic Dynamics & Control*, 143, Article 104505.
- Chib, S. (1993). Bayes regression with autoregressive errors: A Gibbs sampling approach. *Journal of Econometrics*, 58(3), 275–294.
- Chib, S. (1996). Calculating posterior distributions and modal estimates in Markov mixture models. *Journal of Econometrics*, 75(1), 79–97.
- Chib, S. (2011). Introduction to simulation and MCMC methods. In J. Geweke, G. Koop, & H. van Dijk (Eds.), *The Oxford handbook of Bayesian econometrics* (pp. 183–217). OUP.
- Chib, S., & Greenberg, E. (1994). Bayes inference for regression models with ARMA(p,q) errors. *Journal of Econometrics*, 64(1–2), 183–206.
- Chib, S., & Greenberg, E. (1995). Understanding the Metropolis–Hastings algorithm. *American Statistician*, 49(4), 327–335.
- Chib, S., & Greenberg, E. (1996). Markov chain Monte Carlo simulation methods in econometrics. *Economic Theory*, 12(3), 409–431.
- Chib, S., Nadari, F., & Shephard, N. (2002). Markov chain Monte Carlo methods for stochastic volatility models. *Journal of Econometrics*, 108, 281–316.
- Chib, S., Nardari, F., & Shephard, N. (2006). Analysis of high dimensional multivariate stochastic volatility models. *Journal of Econometrics*, 134(2), 341–371.
- Chib, S., Omori, Y., & Asai, M. (2009). Multivariate stochastic volatility. In *Handbook of financial time series* (pp. 365–400). Springer.
- Chipman, H. A., George, E. I., & McCulloch, R. E. (2010). BART: Bayesian additive regression trees. *The Annals of Applied Statistics*, 4(1), 266–298.
- Chiu, C.-W. J., Mumtaz, H., & Pinter, G. (2017). Forecasting with VAR models: Fat tails and stochastic volatility. *International Journal of Forecasting*, 33(4), 1124–1143.
- Chopin, N., & Papaspiliopoulos (2020). *An introduction to sequential Monte Carlo*. Springer.
- Clark, T. E. (2011). Real-time density forecasts from Bayesian vector autoregressions with stochastic volatility. *Journal of Business & Economic Statistics*, 29(3), 327–341.
- Clark, T. E., Huber, F., Koop, G., & Marcellino, M. (2022a). Forecasting US inflation using Bayesian nonparametric models. arXiv preprint arXiv:2202.13793.
- Clark, T., Huber, F., Koop, G., Marcellino, M., & Pfarrhofer, M. (2022b). Tail forecasting with multivariate Bayesian additive regression trees. *International Economic Review*, (Forthcoming).
- Clark, T. E., & Ravazzolo, F. (2015). Macroeconomic forecasting performance under alternative specifications of time-varying volatility. *Journal of Applied Econometrics*, 30(4), 551–575.
- Coelho, C., Stephenson, D., Doblas-Reyes, F., Balmaseda, M., Guetter, A., & Van Oldenborgh, G. (2006). A Bayesian approach for multi-model downscaling: Seasonal forecasting of regional rainfall and river flows in South America. *Meteorological Applications*, 13(1), 73–82.
- Cottet, R., & Smith, M. (2003). Bayesian modeling and forecasting of intraday electricity load. *Journal of the American Statistical Association*, 98(464), 839–849.
- Craiu, R. V., & Meng, X.-L. (2005). Multiprocess parallel antithetic coupling for backward and forward Markov chain Monte Carlo. *The Annals of Statistics*, 33(2), 661–697.
- Creel, M., & Kristensen, D. (2015). ABC of SV: Limited information likelihood inference in stochastic volatility jump-diffusion models. *Journal of Empirical Finance*, 31, 85–108.
- Cross, J. L., Hou, C., & Poon, A. (2020). Macroeconomic forecasting with large Bayesian VARs: Global-local priors and the illusion of sparsity. *International Journal of Forecasting*, 36(3), 899–915.
- Da Silva, F. L., Oliveira, F. L. C., & Souza, R. C. (2019). A bottom-up Bayesian extension for long term electricity consumption forecasting. *Energy*, 167, 198–210.
- D'Agostino, A., Gambetti, L., & Giannone, D. (2013). Macroeconomic forecasting and structural change. *Journal of Applied Econometrics*, 28(1), 82–101.
- Danaher, P. J., Danaher, T. S., Smith, M. S., & Loaiza-Maya, R. (2020). Advertising effectiveness for multiple retailer-brands in a multimedia and multichannel environment. *Journal of Marketing Research*, 57(3), 445–467.
- Davis, P., & Rabinowitz, P. (1975). *Numerical methods of integration*. New York: Academic Press.
- Dawid, A. P. (1982). The well-calibrated Bayesian. *Journal of the American Statistical Association*, 77(379), 605–610.
- Dawid, A. P. (1985). Calibration-based empirical probability. *The Annals of Statistics*, 13(4), 1251–1274.
- Del Negro, M., Hasegawa, R. B., & Schorfheide, F. (2016). Dynamic prediction pools: An investigation of financial frictions and forecasting performance. *Journal of Econometrics*, 192(2), 391–405, *Innovations in Multiple Time Series Analysis*.
- Delatola, E.-I., & Griffin, J. E. (2013). A Bayesian semiparametric model for volatility with a leverage effect. *Computational Statistics & Data Analysis*, 60, 97–110.
- Deligiannidis, G., Doucet, A., & Pitt, M. K. (2018). The correlated pseudomarginal method. *Journal of the Royal Statistical Society Series B*, 80(5), 839–870.
- Dieppe, A., van Roye, B., & Legrand, R. (2016). *The BEAR toolbox: European Central Bank Working Paper 1934*.
- Doan, T., Litterman, R., & Sims, C. (1984). Forecasting and conditional projection using realistic prior distributions. *Econometric Reviews*, 3(1), 1–100.
- Dongarra, J., & Sullivan, F. (2000). Guest editors' introduction: The top 10 algorithms. *Computing in Science & Engineering*, 2(1), 22–23.
- Douc, R., & Robert, C. P. (2011). A vanilla Rao–Blackwellization of Metropolis–Hastings algorithms. *The Annals of Statistics*, 39(1), 261–277.
- Doucet, A., Pitt, M. K., Deligiannidis, G., & Kohn, R. (2015). Efficient implementation of Markov chain Monte Carlo when using an unbiased likelihood estimator. *Biometrika*, 102(2), 295–313.
- Du, P. (2018). Ensemble machine learning-based wind forecasting to combine NWP output with data from weather station. *IEEE Transactions on Sustainable Energy*, 10(4), 2133–2141.
- Dufays, A. (2016). Infinite-state Markov-switching for dynamic volatility. *Journal of Financial Econometrics*, 14(2), 418–460.
- Dunson, D., & Johndrow, J. (2020). The Hastings algorithm at fifty. *Biometrika*, 107(1), 1–23.
- Eraker, B. (2001). MCMC analysis of diffusion models with application to finance. *Journal of Business & Economic Statistics*, 19(2), 177–191.



- Eraker, B. (2004). Do stock prices and volatility jump? Reconciling evidence from spot and option prices. *The Journal of Finance*, 59(3), 1367–1403.
- Eraker, B., Johannes, M., & Polson, N. (2003). The impact of jumps in volatility and returns. *The Journal of Finance*, 58(3), 1269–1300.
- Fan, Y., & Sisson, S. A. (2011). Reversible jump MCMC. In S. Brooks, A. Gelman, G. Jones, & X.-L. Meng (Eds.), *Handbook of Markov chain Monte Carlo* (pp. 67–92). Chapman & Hall/CRC.
- Fasano, A., & Durante, D. (2022). A class of conjugate priors for multinomial probit models which includes the multivariate normal one. *Journal of Machine Learning Research*, 23(30), 1–26.
- Fearnhead, P. (2011). MCMC for state-space models. In S. Brooks, A. Gelman, G. Jones, & X.-L. Meng (Eds.), *Handbook of Markov chain Monte Carlo* (pp. 513–529). Chapman & Hall/CRC.
- Fileccia, G., & Sgarra, C. (2018). A particle filtering approach to oil futures price calibration and forecasting. *Journal of Commodity Markets*, 9, 21–34.
- Flury, T., & Shephard, N. (2011). Bayesian inference based only on simulated likelihood: Particle filter analysis of dynamic economic models. *Economic Theory*, 27(5), 933–956.
- Forbes, C. S., Martin, G. M., & Wright, J. (2007). Inference for a class of stochastic volatility models using option and spot prices: Application of a bivariate Kalman filter. *Econometric Reviews*, 26(2–4), 387–418.
- Frazier, D. T., Loaiza-Maya, R., & Martin, G. M. (2022). Variational Bayes in state space models: Inferential and predictive accuracy. *Journal of Computational and Graphical Statistics*, <https://doi.org/10.1080/10618600.2022.2134875>, (Forthcoming).
- Frazier, D. T., Loaiza-Maya, R., Martin, G. M., & Koo, B. (2022). Loss-based variational Bayes prediction. arXiv preprint arXiv:2104.14054.
- Frazier, D. T., Maneesoonthorn, W., Martin, G. M., & McCabe, B. P. (2019). Approximate Bayesian forecasting. *International Journal of Forecasting*, 35(2), 521–539.
- Frühwirth-Schnatter, S. (1994). Data augmentation and dynamic linear models. *Journal of Time Series Analysis*, 15(2), 183–202.
- Frühwirth-Schnatter, S. (2004). Efficient Bayesian parameter estimation. In J. Harvey, S. Koopman, & N. Shephard (Eds.), *State space and unobserved component Models: Theory and applications* (pp. 123–151). OUP, CUP.
- Frühwirth-Schnatter, S., & Wagner, H. (2010). Stochastic model specification search for Gaussian and partial non-Gaussian state space models. *Journal of Econometrics*, 154(1), 85–100.
- Fulop, A., & Li, J. (2019). Bayesian estimation of dynamic asset pricing models with informative observations. *Journal of Econometrics*, 209(1), 114–138.
- Gallant, A. R., & Tauchen, G. (1996). Which moments to match? *Economic Theory*, 12(4), 657–681.
- Gefang, D., Koop, G., & Poon, A. (2022). Forecasting using variational Bayesian inference in large vector autoregressions with hierarchical shrinkage. *International Journal of Forecasting*, 39(1), 346–363.
- Gelfand, A., & Smith, A. (1990). Sampling based approaches to calculating marginal densities. *Journal of the American Statistical Association*, 85(410), 398–409.
- Geman, S., & Geman, D. (1984). Stochastic relaxation, Gibbs distributions and the Bayesian restoration of images. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 6, 721–741.
- George, E. (2000). The variable selection problem. *Journal of the American Statistical Association*, 95(452), 1304–1308.
- George, E. I., Sun, D., & Ni, S. (2008). Bayesian stochastic search for VAR model restrictions. *Journal of Econometrics*, 142(1), 553–580.
- Geraci, M. V., & Gnabo, J.-Y. (2018). Measuring interconnectedness between financial institutions with Bayesian time-varying vector autoregressions. *Journal of Financial and Quantitative Analysis*, 53(3), 1371–1390.
- Geweke, J. (1989). Bayesian inference in econometric models using Monte Carlo integration. *Econometrica*, 57(6), 1317–1340.
- Geweke, J. (2005). *Contemporary Bayesian econometrics and statistics*, vol. 537. John Wiley & Sons.
- Geweke, J., & Amisano, G. (2011). Optimal prediction pools. *Journal of Econometrics*, 164(1), 130–141.
- Geweke, J., & Whiteman, C. (2006). Bayesian forecasting. In *Handbook of economic forecasting*, vol. 1 (pp. 3–80). Elsevier.
- Geyer, C. J. (2011). Introduction to Markov chain Monte Carlo. In S. Brooks, A. Gelman, G. Jones, & X.-L. Meng (Eds.), *Handbook of Markov chain Monte Carlo* (pp. 3–48). Chapman & Hall/CRC.
- Ghayekhloo, M., Azimi, R., Ghofrani, M., Menhaj, M., & Shekari, E. (2019). A combination approach based on a novel data clustering method and Bayesian recurrent neural network for day-ahead price forecasting of electricity markets. *Electric Power Systems Research*, 168, 184–199.
- Gianfreda, A., Ravazzolo, F., & Rossini, L. (2020). Comparing the forecasting performances of linear models for electricity prices with high RES penetration. *International Journal of Forecasting*, 36(3), 974–986.
- Giannone, D., Lenza, M., & Primiceri, G. E. (2015). Prior selection for vector autoregressions. *The Review of Economics and Statistics*, 97(2), 436–451.
- Giebel, G., & Kariniotakis, G. (2017). Wind power forecasting—A review of the state of the art. In *Renewable Energy Forecasting: From models to applications* (pp. 59–109). Woodhead Publishing.
- Gilanifar, M., Wang, H., Ozguven, E. E., Zhou, Y., & Arghandeh, R. (2019). Bayesian spatiotemporal Gaussian process for short-term load forecasting using combined transportation and electricity data. *ACM Transactions on Cyber-Physical Systems*, 4(1), 1–25.
- Gilbride, T. J., & Allenby, G. M. (2004). A choice model with conjunctive, disjunctive, and compensatory screening rules. *Marketing Science*, 23(3), 391–406.
- Giordani, P., Pitt, M., & Kohn, R. (2011). Bayesian inference for time series state space models. In J. Geweke, G. Koop, & H. van Dijk (Eds.), *The Oxford handbook of Bayesian econometrics* (pp. 61–124). OUP.
- Girolami, M., & Rogers, S. (2006). Variational Bayesian multinomial probit regression with Gaussian process priors. *Neural Computation*, 18(8), 1790–1817.
- Giummolè, F., Mameli, V., Ruli, E., & Ventura, L. (2017). Objective Bayesian inference with proper scoring rules. *TEST*, 28(3), 1–28.
- Glynn, P. W., & Rhee, C.-H. (2014). Exact estimation for Markov chain equilibrium expectations. *Journal of Applied Probability*, 51(A), 377–389.
- Gneiting, T., Balabdaoui, F., & Raftery, A. E. (2007). Probabilistic forecasts, calibration and sharpness. *Journal of the Royal Statistical Society. Series B. Statistical Methodology*, 69(2), 243–268.
- Gneiting, T., & Raftery, A. E. (2007). Strictly proper scoring rules, prediction, and estimation. *Journal of the American Statistical Association*, 102(477), 359–378.
- Gneiting, T., & Ranjan, R. (2013). Combining predictive distributions. *Electronic Journal of Statistics*, 7, 1747–1782.
- Gonzato, L., & Sgarra, C. (2021). Self-exciting jumps in the oil market: Bayesian estimation and dynamic hedging. *Energy Economics*, 99, Article 105279.
- Gordon, N., Salmond, J., & Smith, A. (1993). A novel approach to non-linear/non-Gaussian Bayesian state estimation. *IEEE Proceedings on Radar and Signal Processing*, 140(2), 107–113.
- Granger, C. W., et al. (1986). Forecasting accuracy of alternative techniques: A comparison of US macroeconomic forecasts: Comment. *Journal of Business & Economic Statistics*, 4(1), 16–17.
- Green, P. (1995). Reversible jump MCMC computation and Bayesian model determination. *Biometrika*, 82(4), 711–732.
- Green, P. J. (2003). Trans-dimensional Markov chain Monte Carlo. In P. J. Green, H. N. L., & R. S. (Eds.), *Oxford statistical science series*, vol. 27 (pp. 179–198).
- Green, P., Latuszynski, K., Pereyra, M., & Robert, C. (2015). Bayesian computation: A summary of the current state, and samples backwards and forwards. *Statistics and Computing*, 25, 835–862.
- Griffin, J. E., & Steel, M. F. (2011). Stick-breaking autoregressive processes. *Journal of Econometrics*, 162(2), 383–396.
- Grillone, B., Mor, G., Danov, S., Cipriano, J., Lazzari, F., & Sumper, A. (2021). Baseline energy use modeling and characterization in tertiary buildings using an interpretable Bayesian linear regression methodology. *Energies*, 14(17), 5556.
- Guedj, B. (2019). A primer on PAC-Bayesian learning. arXiv preprint arXiv:1901.05353.
- Gunawan, D., Kohn, R., & Nott, D. (2021). Variational Bayes approximation of factor stochastic volatility models. *International Journal of Forecasting*, 37(4), 1355–1375.
- Gunel, I. (1987). Forecasting system energy demand. *Journal of Forecasting*, 6(2), 137–156.

- Hafner, C. M., & Herwartz, H. (2001). Option pricing under linear autoregressive dynamics, heteroskedasticity, and conditional leptokurtosis. *Journal of Empirical Finance*, 8(1), 1–34.
- Hall, S. G., & Mitchell, J. (2007). Combining density forecasts. *International Journal of Forecasting*, 23(1), 1–13.
- Hammersley, J., & Handscomb, D. (1964). *Monte Carlo methods*. New York: John Wiley.
- Harvey, A. (1981). *The econometric analysis of time series*. John Wiley.
- Hassan, S., Khosravi, A., & Jaafar, J. (2015). Examining performance of aggregation algorithms for neural network-based electricity demand forecasting. *International Journal of Electrical Power & Energy Systems*, 64, 1098–1105.
- Hastings, W. (1970). Monte Carlo sampling methods using Markov chains and their application. *Biometrika*, 57(1), 97–109.
- Hauzenberger, N., Huber, F., Koop, G., & Onorante, L. (2022). Fast and flexible Bayesian inference in time-varying parameter regression models. *Journal of Business & Economic Statistics*, 40(4), 1904–1918.
- Hauzenberger, N., Huber, F., & Onorante, L. (2021). Combining shrinkage and sparsity in conjugate vector autoregressive models. *Journal of Applied Econometrics*, 36(3), 304–327.
- Hippert, H. S., & Taylor, J. W. (2010). An evaluation of Bayesian techniques for controlling model complexity and selecting inputs in a neural network for short-term load forecasting. *Neural Networks*, 23(3), 386–395.
- Hoffman, M. D., & Gelman, A. (2014). The no-u-turn sampler: Adaptively setting path lengths in Hamiltonian Monte Carlo. *Journal of Machine Learning Research*, 15(1), 1593–1623.
- Holmes, C., & Walker, S. (2017). Assigning a value to a power likelihood in a general Bayesian model. *Biometrika*, 104(2), 497–503.
- Huber, M. L. (2016). *Perfect simulation*. Chapman & Hall/CRC.
- Huber, F., & Feldkircher, M. (2019). Adaptive shrinkage in Bayesian vector autoregressive models. *Journal of Business & Economic Statistics*, 37(1), 27–39.
- Huber, F., Koop, G., & Onorante, L. (2021). Inducing sparsity and shrinkage in time-varying parameter models. *Journal of Business & Economic Statistics*, 39(3), 669–683.
- Huber, F., Koop, G., Onorante, L., Pfarrhofer, M., & Schreiner, J. (2023). Nowcasting in a pandemic using non-parametric mixed frequency VARs. *Journal of Econometrics*, 232(1), 52–69.
- Huber, F., & Pfarrhofer, M. (2021). Dynamic shrinkage in time-varying parameter stochastic volatility in mean models. *Journal of Applied Econometrics*, 36(2), 262–270.
- Jacob, P. E., O'Leary, J., & Atchadé, Y. F. (2020). Unbiased Markov chain Monte Carlo methods with couplings. *Journal of the Royal Statistical Society Series B*, 82, 1–32, With discussion.
- Jacob, P., Robert, C., & Smith, M. (2011). Using parallel computation to improve independent Metropolis–Hastings based estimation. *Journal of Computational and Graphical Statistics*, 20(3), 616–635.
- Jacquier, E., & Polson, N. (2011). Bayesian methods in finance. In J. Geweke, G. Koop, & H. van Dijk (Eds.), *The Oxford handbook of Bayesian econometrics* (pp. 439–512). OUP.
- Jacquier, R., Polson, N. G., & Rossi, P. E. (1994). Bayesian analysis of stochastic volatility models. *Journal of Business & Economic Statistics*, 12(4), 371–389, With discussion.
- Jahan, F., Ullah, I., & Mengersen, K. (2020). A review of Bayesian statistical approaches for big data. In K. Mengersen, P. Pudlo, & C. Robert (Eds.), *Case studies in applied Bayesian science* (pp. 17–44). Springer.
- Jensen, M. J., & Maheu, J. M. (2010). Bayesian semiparametric stochastic volatility modeling. *Journal of Econometrics*, 157(2), 306–316.
- Jensen, M. J., & Maheu, J. M. (2013). Bayesian semiparametric multivariate GARCH modeling. *Journal of Econometrics*, 176(1), 3–17.
- Jensen, M. J., & Maheu, J. M. (2014). Estimating a semiparametric asymmetric stochastic volatility model with a Dirichlet process mixture. *Journal of Econometrics*, 178(3), 523–538.
- Jin, X., & Maheu, J. M. (2013). Modeling realized covariances and returns. *Journal of Financial Econometrics*, 11(2), 335–369.
- Jin, X., & Maheu, J. M. (2016). Bayesian semiparametric modeling of realized covariance matrices. *Journal of Econometrics*, 192(1), 19–39.
- Jin, X., Maheu, J. M., & Yang, Q. (2019). Bayesian parametric and semi-parametric factor models for large realized covariance matrices. *Journal of Applied Econometrics*, 34(5), 641–660.
- Jin, X., Maheu, J. M., & Yang, Q. (2022). Infinite Markov pooling of predictive distributions. *Journal of Econometrics*, 228(2), 302–321.
- Johannes, M., & Polson, N. (2010). CHAPTER 13 – MCMC methods for continuous-time financial econometrics. In Y. Ait-Sahalia, & L. P. Hansen (Eds.), *Handbooks in finance: vol. 2, Handbook of financial econometrics: applications* (pp. 1–72). San Diego: Elsevier.
- Johannes, M. S., Polson, N. G., & Stroud, J. R. (2009). Optimal filtering of jump diffusions: Extracting latent states from asset prices. *The Review of Financial Studies*, 22(7), 2759–2799.
- Johnrow, J. E., Smith, A., Pillai, N., & Dunson, D. B. (2019). MCMC for imbalanced categorical data. *Journal of the American Statistical Association*, 114(527), 1394–1403.
- Johnson, M. C. (2017). *Bayesian predictive synthesis: Forecast calibration and combination* (Ph.D. thesis), Duke University.
- Joutz, F. L., Maddala, G. S., & Trost, R. P. (1995). An integrated Bayesian vector auto regression and error correction model for forecasting electricity consumption and prices. *Journal of Forecasting*, 14(3), 287–310.
- Kabisa, S., Dunson, D. B., & Morris, J. S. (2016). Online variational Bayes inference for high-dimensional correlated data. *Journal of Computational and Graphical Statistics*, 25(2), 426–444.
- Kalli, M., & Griffin, J. (2015). Flexible modeling of dependence in volatility processes. *Journal of Business & Economic Statistics*, 33(1), 102–113.
- Kalli, M., & Griffin, J. E. (2018). Bayesian nonparametric vector autoregressive models. *Journal of Econometrics*, 203(2), 267–282.
- Kalli, M., Griffin, J. E., & Walker, S. G. (2011). Slice sampling mixture models. *Statistics and Computing*, 21(1), 93–105.
- Karmakar, B., Kwon, O., Mukherjee, G., & Siddharth, S. (2021). *Understanding early adoption of hybrid cars via a new multinomial probit model with multiple network weights: Technical report*.
- Kastner, G., Frühwirth-Schnatter, S., & Lopes, H. F. (2017). Efficient Bayesian inference for multivariate factor stochastic volatility models. *Journal of Computational and Graphical Statistics*, 26(4), 905–917.
- Kastner, G., & Huber, F. (2021). Sparse Bayesian vector autoregressions in huge dimensions. *Journal of Forecasting*, 39(7), 1142–1165.
- Kim, S., Shephard, N., & Chib, S. (1998). Stochastic volatility: Likelihood inference and comparison with ARCH models. *Review of Economic Studies*, 65(3), 361–393.
- Kloek, T., & van Dijk, H. K. (1978). Bayesian estimates of equation system parameters: An application of integration by Monte Carlo. *Econometrica*, 46(1), 1–19.
- Kon Kam King, G., Canale, A., & Ruggiero, M. (2019). Bayesian functional forecasting with locally-autoregressive dependent processes. *Bayesian Analysis*, 14(4), 1121–1141.
- Koop, G. M. (2003). *Bayesian econometrics*. John Wiley & Sons Inc..
- Koop, G. (2013a). Forecasting with medium and large Bayesian VARs. *Journal of Applied Econometrics*, 28(2), 177–203.
- Koop, G. M. (2013b). Forecasting with medium and large Bayesian VARs. *Journal of Applied Econometrics*, 28(2), 177–203.
- Koop, G., & Korobilis, D. (2010). Bayesian multivariate time series methods for empirical macroeconomics. *Foundations and Trends® in Econometrics*, 3(4), 267–358.
- Koop, G., & Korobilis, D. (2012). Forecasting inflation using dynamic model averaging. *International Economic Review*, 53(3), 867–886.
- Koop, G., & Korobilis, D. (2013). Large time-varying parameter VARs. *Journal of Econometrics*, 177(2), 185–198.
- Koop, G., & Korobilis, D. (2023). Bayesian dynamic variable selection in high dimensions. *International Economic Review*, (Forthcoming).
- Koop, G., McIntyre, S., Mitchell, J., & Poon, A. (2020). Regional output growth in the United Kingdom: More timely and higher frequency estimates from 1970. *Journal of Applied Econometrics*, 35(2), 176–197.
- Korobilis, D. (2013). VAR forecasting using Bayesian variable selection. *Journal of Applied Econometrics*, 28(2), 204–230.
- Kostrzewski, M., & Kostrzewska, J. (2019). Probabilistic electricity price forecasting with Bayesian stochastic volatility models. *Energy Economics*, 80, 610–620.
- Lahiri, K., & Gao, J. (2002). Bayesian analysis of nested logit model by Markov chain Monte Carlo. *Journal of Econometrics*, 111(1), 103–133.
- Laplace, P. (1774). Mémoire sur la probabilité des causes par les événements. *Mémoires De l'académie Royale Des Sciences Présentés Par Divers Savants*, 6, 621–656.

- Launay, T., Philippe, A., & Lamarche, S. (2015). Construction of an informative hierarchical prior for a small sample with the help of historical data and application to electricity load forecasting. *Test*, 24(2), 361–385.
- Lenza, M., & Primiceri, G. E. (2022). How to estimate a vector autoregression after March 2020. *Journal of Applied Econometrics*, 37(4), 688–699.
- Li, C. (2022). *A multivariate GARCH model with an infinite hidden Markov mixture*: MPRA Paper No. 112792.
- Lim, G.-C., Martin, G. M., & Martin, V. L. (2005). Parametric pricing of higher order moments in S&P500 options. *Journal of Applied Econometrics*, 20(3), 377–404.
- Lindberg, K., Seljom, P., Madsen, H., Fischer, D., & Korpås, M. (2019). Long-term electricity load forecasting: Current and future trends. *Utilities Policy*, 58, 102–119.
- Lintusaari, J., Gutmann, M. U., Dutta, R., Kaski, S., & Corander, J. (2017). Fundamentals and recent developments in approximate Bayesian computation. *Systematic Biology*, 66(1), e66–e82.
- Liu, J., & Maheu, J. M. (2018). Improving Markov switching models using realized variance. *Journal of Applied Econometrics*, 33(3), 297–318.
- Liu, J., Wong, W., & Kong, A. (1994). Covariance structure of the Gibbs sampler with application to the comparison of estimators and augmentation schemes. *Biometrika*, 81(1), 27–40.
- Llorente, F., Martino, L., Delgado, D., & Lopez-Santiago, J. (2021). Marginal likelihood computation for model selection and hypothesis testing: An extensive review. *Siam Review*, 65(1), <https://doi.org/10.1137/20M1310849>.
- Loaiza-Maya, R., Martin, G. M., & Frazier, D. T. (2021). Focused Bayesian prediction. *Journal of Applied Econometrics*, 36(5), 517–543.
- Loaiza-Maya, R., & Nibbering, D. (2022a). Fast variational Bayes methods for multinomial probit models. *Journal of Business & Economic Statistics*, <https://doi.org/10.1080/07350015.2022.2139267>, (Forthcoming).
- Loaiza-Maya, R., & Nibbering, D. (2022b). Scalable Bayesian estimation in the multinomial probit model. *Journal of Business & Economic Statistics*, 40(4), 1678–1690.
- Loaiza-Maya, R., Smith, M. S., Nott, D. J., & Danaher, P. J. (2022). Fast and accurate variational inference for models with many latent variables. *Journal of Econometrics*, 230(2), 339–362.
- Lyddon, S., Holmes, C., & Walker, S. (2019). General Bayesian updating and the loss-likelihood bootstrap. *Biometrika*, 106(2), 465–478.
- Madigan, D., & Raftery, A. (1995). Bayesian graphical models for discrete data. *International Statistical Review*, 63(2), 215–232.
- Maneesoonthorn, W., Forbes, C. S., & Martin, G. M. (2017). Inference on self-exciting jumps in prices and volatility using high-frequency measures. *Journal of Applied Econometrics*, 32(3), 504–532.
- Maneesoonthorn, W., Martin, G. M., Forbes, C. S., & Grose, S. D. (2012). Probabilistic forecasts of volatility and its risk premia. *Journal of Econometrics*, 171(2), 217–236.
- Marin, J.-M., Mengersen, K., & Robert, C. (2005). Bayesian modelling and inference on mixtures of distributions. In C. Rao, & D. Dey (Eds.), *Handbook of statistics* (pp. 459–507). New York: Elsevier.
- Marin, J., Pudlo, P., Robert, C., & Ryder, R. (2011). Approximate Bayesian computational methods. *Statistics and Computing*, 21(2), 279–291.
- Martin, G. M., Frazier, D. T., & Robert, C. P. (2023a). Approximating Bayes in the 21st century. *Statistical Science*, 38, <https://doi.org/10.1214/22-STS875>.
- Martin, G. M., Frazier, D. T., & Robert, C. P. (2023b). Computing Bayes: From then 'Til now. *Statistical Science*, 38, <https://doi.org/10.1214/22-STS876>.
- Martin, G. M., Loaiza-Maya, R., Maneesoonthorn, W., Frazier, D. T., & Ramirez-Hassan, A. (2022). Optimal probabilistic forecasts: When do they work? *International Journal of Forecasting*, 38(1), 384–406.
- Martin, G. M., McCabe, B. P., Frazier, D. T., Maneesoonthorn, W., & Robert, C. P. (2019). Auxiliary likelihood-based approximate Bayesian computation in state space models. *Journal of Computational and Graphical Statistics*, 28(3), 508–522.
- Martino, S., & Riebler, A. (2019). Integrated Nested Laplace Approximations (INLA). arXiv preprint [arXiv:1907.01248](https://arxiv.org/abs/1907.01248).
- McAlinn, K., Aastveit, K. A., Nakajima, J., & West, M. (2020). Multivariate Bayesian predictive synthesis in macroeconomic forecasting. *Journal of the American Statistical Association*, 115(531), 1092–1110.
- McAlinn, K., & West, M. (2019). Dynamic Bayesian predictive synthesis in time series forecasting. *Journal of Econometrics*, 210(1), 155–169.
- McCormick, T. H., Raftery, A. E., Madigan, D., & Burd, R. S. (2012). Dynamic logistic regression and dynamic model averaging for binary classification. *Biometrics*, 68(1), 23–30.
- McCracken, M., Owyang, M., & Sekhposyan, T. (2021). Real-time forecasting and scenario analysis with a large mixed-frequency Bayesian VAR. *International Journal of Central Banking*, 18(5), 327–367.
- McCulloch, R. E., Polson, N. G., & Rossi, P. E. (2000). A Bayesian analysis of the multinomial probit model with fully identified parameters. *Journal of Econometrics*, 99(1), 173–193.
- McCulloch, R., & Rossi, P. E. (1994). An exact likelihood analysis of the multinomial probit model. *Journal of Econometrics*, 64(1–2), 207–240.
- McCulloch, R. E., & Tsay, R. S. (1994). Bayesian analysis of autoregressive time series via the Gibbs sampler. *Journal of Time Series Analysis*, 15(2), 235–250.
- McFadden, D. (1989). A method of simulated moments for estimation of discrete response models without numerical integration. *Econometrica*, 57(5), 995–1026.
- Metropolis, N., Rosenbluth, A. W., Rosenbluth, M. N., Teller, A. H., & Teller, E. (1953). Equations of state calculations by fast computing machines. *Journal of Chemical Physics*, 21, 1087–1092.
- Metropolis, N., & Ulam, S. (1949). The Monte Carlo method. *Journal of the American Statistical Association*, 44(247), 335–341.
- Mishkin, E. (2021). Gender and sibling dynamics in the intergenerational transmission of entrepreneurship. *Management Science*, 67(10), 6116–6135.
- Mittal, V., Han, K., Lee, J.-Y., & Sridhar, S. (2021). Improving business-to-business customer satisfaction programs: Assessment of asymmetry, heterogeneity, and financial impact. *Journal of Marketing Research*, 58(4), 615–643.
- Miyazaki, K., Hoshino, T., & Böckenholt, U. (2021). Dynamic two stage modeling for category-level and brand-level purchases using potential outcome approach with Bayes inference. *Journal of Business & Economic Statistics*, 39(3), 622–635.
- Mori, H., & Nakano, K. (2014). Application of Gaussian process to locational marginal pricing forecasting. *Procedia Computer Science*, 36, 220–226.
- Müller, G., & Uhl, S. (2021). Estimation of time-varying autoregressive stochastic volatility models with stable innovations. *Statistics and Computing*, 31(3), 1–19.
- Naesseth, C. A., Lindsten, F., Schön, T. B., et al. (2019). Elements of sequential Monte Carlo. *Foundations and Trends in Machine Learning*, 12(3), 307–392.
- Nakajima, J. (2017). Bayesian analysis of multivariate stochastic volatility with skew return distribution. *Econometric Reviews*, 36(5), 546–562.
- Nateghi, R., Guikema, S. D., & Quiring, S. M. (2011). Comparison and validation of statistical methods for predicting power outage durations in the event of hurricanes. *Risk Analysis: An International Journal*, 31(12), 1897–1906.
- Naylor, J., & Smith, A. (1982). Application of a method for the efficient computation of posterior distributions. *Applied Statistics*, 31(3), 214–225.
- Neal, R. (2011a). MCMC using ensembles of states for problems with fast and slow variables such as Gaussian process regression. arXiv preprint [arXiv:1101.0387](https://arxiv.org/abs/1101.0387).
- Neal, R. (2011b). MCMC using Hamiltonian dynamics. In S. Brooks, A. Gelman, G. Jones, & X.-L. Meng (Eds.), *Handbook of Markov chain Monte Carlo* (pp. 113–162). Chapman & Hall/CRC.
- Neiswanger, W., Wang, C., & Xing, E. (2013). Asymptotically exact, embarrassingly parallel MCMC. arXiv preprint [arXiv:1311.4780](https://arxiv.org/abs/1311.4780).
- Nonejad, N. (2021). An overview of dynamic model averaging techniques in time-series econometrics. *Journal of Economic Surveys*, 35(2), 566–614.
- Nott, D. J., & Kohn, R. (2005). Adaptive sampling for Bayesian variable selection. *Biometrika*, 92(4), 747–763.
- Nowotarski, J., Raviv, E., Trück, S., & Weron, R. (2014). An empirical comparison of alternative schemes for combining electricity spot price forecasts. *Energy Economics*, 46, 395–412.
- Nowotarski, J., & Weron, R. (2018). Recent advances in electricity price forecasting: A review of probabilistic forecasting. *Renewable and Sustainable Energy Reviews*, 81(1), 1548–1568.



- Ohtsuka, Y., Oga, T., & Kakamu, K. (2010). Forecasting electricity demand in Japan: A Bayesian spatial autoregressive ARMA approach. *Computational Statistics & Data Analysis*, 54(11), 2721–2735.
- Omori, Y., Chib, S., Shephard, N., & Nakajima, J. (2007). Stochastic volatility with leverage: Fast and efficient likelihood inference. *Journal of Econometrics*, 140(2), 425–449.
- Opschoor, A., Van Dijk, D., & van der Wel, M. (2017). Combining density forecasts using focused scoring rules. *Journal of Applied Econometrics*, 32(7), 1298–1313.
- Ormerod, J. T., & Wand, M. P. (2010). Explaining variational approximations. *American Statistician*, 64(2), 140–153.
- Owen, A. B. (2017). Statistically efficient thinning of a Markov chain sampler. *Journal of Computational and Graphical Statistics*, 26(3), 738–744.
- Paleti, R. (2018). Generalized multinomial probit model: Accommodating constrained random parameters. *Transportation Research, Part B (Methodological)*, 118, 248–262.
- Panagiotelis, A., & Smith, M. (2008). Bayesian density forecasting of intraday electricity prices using multivariate skew t distributions. *International Journal of Forecasting*, 24(4), 710–727.
- Pesonen, H., Simola, U., Köhn-Luque, A., Vuollekoski, H., Lai, X., Frigessi, A., et al. (2022). ABC of the future. *International Statistical Review*, <https://doi.org/10.1111/insr.12522>.
- Peters, G. W., Sisson, S. A., & Fan, Y. (2012). Likelihood-free Bayesian inference for  $\alpha$ -stable models. *Computational Statistics & Data Analysis*, 56(11), 3743–3756.
- Pettenuzzo, D., & Ravazzolo, F. (2016). Optimal portfolio choice under decision-based model combinations. *Journal of Applied Econometrics*, 31(7), 1312–1332.
- Pezzulli, S., Frederic, P., Majithia, S., Sabbagh, S., Black, E., Sutton, R., et al. (2006). The seasonal forecast of electricity demand: A hierarchical Bayesian model with climatological weather generator. *Applied Stochastic Models in Business and Industry*, 22(2), 113–125.
- Pitt, M. K., dos Santos Silva, R., Giordani, P., & Kohn, R. (2012). On some properties of Markov chain Monte Carlo simulation methods based on the particle filter. *Journal of Econometrics*, 171(2), 134–151.
- Poirier, D. J. (1996). A Bayesian analysis of nested logit models. *Journal of Econometrics*, 75(1), 163–181.
- Polson, N. G., Carlin, B. P., & Stoffer, D. S. (1992). A Monte Carlo approach to nonnormal and nonlinear state-space modeling. *Journal of the American Statistical Association*, 87(418), 493–500.
- Posch, K., Truden, C., Hungerländer, P., & Pilz, J. (2022). A Bayesian approach for predicting food and beverage sales in staff canteens and restaurants. *International Journal of Forecasting*, 38(1), 321–338.
- Price, L. F., Drovandi, C. C., Lee, A., & Nott, D. J. (2018). Bayesian synthetic likelihood. *Journal of Computational and Graphical Statistics*, 27(1), 1–11.
- Primiceri, G. E. (2005). Time varying structural vector autoregressions and monetary policy. *Review of Economic Studies*, 72(3), 821–852.
- Quiroz, M., Kohn, R., Villani, M., & Tran, M.-N. (2019). Speeding up MCMC by efficient data subsampling. *Journal of the American Statistical Association*, 114(526), 831–843.
- Quiroz, M., Nott, D. J., & Kohn, R. (2022). Gaussian variational approximation for high-dimensional state space models. *Bayesian Analysis*, <https://doi.org/10.1214/22-BA1332>, (Forthcoming).
- Quiroz, M., Tran, M.-N., Villani, M., & Kohn, R. (2018). Speeding up MCMC by delayed acceptance and data subsampling. *Journal of Computational and Graphical Statistics*, 27(1), 12–22.
- Raftery, A. E., Gneiting, T., Balabdaoui, F., & Polakowski, M. (2005). Using Bayesian model averaging to calibrate forecast ensembles. *Monthly Weather Review*, 133(5), 1155–1174.
- Raftery, A., Madigan, D., & Hoeting, J. A. (1997). Bayesian model averaging for linear regression models. *Journal of the American Statistical Association*, 92(437), 179–191.
- Raftery, A., Miroslav, D., & Ettler, P. (2010). Online prediction under model uncertainty via dynamic model averaging: Application to a cold rolling mill. *Technometrics*, 52(1), 52–66.
- Ranjan, R., & Gneiting, T. (2010). Combining probability forecasts. *Journal of the Royal Statistical Society. Series B. Statistical Methodology*, 72(1), 71–91.
- Raviv, E., Bouwman, K. E., & Van Dijk, D. (2015). Forecasting day-ahead electricity prices: Utilizing hourly prices. *Energy Economics*, 50, 227–239.
- Raza, M. Q., Nadarajah, M., & Ekanayake, C. (2017). Demand forecast of PV integrated bioclimatic buildings using ensemble framework. *Applied Energy*, 208, 1626–1638.
- Ritter, C., & Tanner, M. (1992). Facilitating the Gibbs sampler: The Gibbs stopper and the Griddy-Gibbs sampler. *Journal of the American Statistical Association*, 87(419), 861–868.
- Robert, C. (2007). *The Bayesian choice*. New York: Springer-Verlag.
- Robert, C., & Casella, G. (2011). A history of Markov chain Monte Carlo—Subjective recollections from incomplete data. *Statistical Science*, 26(1), 102–115.
- Robert, C. P., Elvira, V., Tawn, N., & Wu, C. (2018). Accelerating MCMC algorithms. *Wiley Interdisciplinary Reviews: Computational Statistics*, 10(5), Article e1435.
- Roberts, G. O., & Rosenthal, J. S. (2009). Examples of adaptive MCMC. *Journal of Computational and Graphical Statistics*, 18(2), 349–367.
- Roberts, G., & Sahu, S. (1997). Updating schemes, covariance structure, blocking and parametrisation for the Gibbs sampler. *Journal of the Royal Statistical Society Series B*, 59(2), 291–317.
- Rossi, P. E., & Allenby, G. M. (2003). Bayesian statistics and marketing. *Marketing Science*, 22(3), 304–328.
- Rossi, P., Allenby, G., & McCulloch, R. (2012). *Wiley series in probability and statistics, Bayesian statistics and marketing*. Wiley.
- Rue, H., Martino, S., & Chopin, N. (2009). Approximate Bayesian inference for latent Gaussian models using integrated nested Laplace approximations. *Journal of the Royal Statistical Society Series B*, 71(2), 319–392.
- Rue, H., Riebler, A., Sørbye, S. H., Illian, J. B., Simpson, D. P., & Lindgren, F. K. (2017). Bayesian computing with INLA: A review. *Annual Review of Statistics and Its Application*, 4(1), 395–421.
- Schorfheide, F., & Song, D. (2015). Real-time forecasting with a mixed-frequency VAR. *Journal of Business & Economic Statistics*, 33(3), 366–380.
- Sethuraman, J. (1994). A constructive definition of Dirichlet priors. *Statistica Sinica*, 4(2), 639–650.
- Shephard, N., & Pitt, M. K. (1997). Likelihood analysis of non-Gaussian measurement times series. *Biometrika*, 84(3), 653–667.
- Shi, S., & Song, Y. (2016). Identifying speculative bubbles using an infinite hidden Markov model. *Journal of Financial Econometrics*, 14(1), 159–184.
- Sisson, S., & Fan, Y. (2011). Likelihood-free Markov chain Monte Carlo. In S. Brooks, A. Gelman, G. Jones, & X.-L. Meng (Eds.), *Handbook of Markov chain Monte Carlo* (pp. 313–333). Chapman & Hall/CRC.
- Sisson, S. A., Fan, Y., & Beaumont, M. (2019). *Handbook of approximate Bayesian computation*. Chapman & Hall/CRC.
- Slouther, J. M., Gneiting, T., & Raftery, A. E. (2010). Probabilistic wind speed forecasting using ensembles and Bayesian model averaging. *Journal of the American Statistical Association*, 105(489), 25–35.
- Smets, F., & Wouters, R. (2007). Shocks and frictions in US business cycles: A Bayesian DSGE approach. *American Economic Review*, 97(3), 586–606.
- Smith, M. (2000). Modeling and short-term forecasting of New South Wales electricity system load. *Journal of Business & Economic Statistics*, 18(4), 465–478.
- Smith, M. S. (2010). Bayesian inference for a periodic stochastic volatility model of intraday electricity prices. In *Statistical modelling and regression structures* (pp. 353–376). Springer.
- Steel, M. F. J. (2020). Model averaging and its use in economics. *Journal of Economic Literature*, 58(3), 644–719.
- Stock, J. H., & Watson, M. (2007). Why has US inflation become harder to forecast? *Journal of Money, Credit and Banking*, 39(1), 3–33.
- Stock, J. H., & Watson, M. (2011). Dynamic factor models. In *Oxford handbooks online*. Oxford University Press.
- Stock, J. H., & Watson, M. W. (2016). Core inflation and trend inflation. *The Review of Economics and Statistics*, 98(4), 770–784.
- Strickland, C. M., Forbes, C. S., & Martin, G. M. (2006). Bayesian analysis of the stochastic conditional duration model. *Computational Statistics & Data Analysis*, 50(9), 2247–2267.
- Strickland, C. M., Martin, G. M., & Forbes, C. S. (2008). Parameterisation and efficient MCMC estimation of non-Gaussian state space models. *Computational Statistics & Data Analysis*, 52(6), 2911–2930.
- Stroud, J. R., Müller, P., & Polson, N. G. (2003). Nonlinear state-space models with state-dependent variances. *Journal of the American Statistical Association*, 98(462), 377–386.



- Sun, P., Kim, I., & Lee, K. (2020). Flexible weighted Dirichlet process mixture modelling and evaluation to address the problem of forecasting return distribution. *Journal of Nonparametric Statistics*, 32(4), 989–1014.
- Sun, M., Zhang, T., Wang, Y., Strbac, G., & Kang, C. (2019). Using Bayesian deep learning to capture uncertainty for residential net load forecasting. *IEEE Transactions on Power Systems*, 35(1), 188–201.
- Syring, N., & Martin, R. (2019). Calibrating general posterior credible regions. *Biometrika*, 106(2), 479–486.
- Tallman, E., & West, M. (2022). Bayesian predictive decision synthesis. arXiv preprint arXiv:2206.03815.
- Tanner, M. A., & Wong, W. (1987). The calculation of posterior distributions by data augmentation. *Journal of the American Statistical Association*, 82(398), 528–550. With discussion.
- Tavaré, S., Balding, D., Griffith, R., & Donnelly, P. (1997). Inferring coalescence times from DNA sequence data. *Genetics*, 145(2), 505–518.
- Teh, Y. W., Jordan, M. I., Beal, M. J., & Blei, D. M. (2006). Hierarchical Dirichlet processes. *Journal of the American Statistical Association*, 101(476), 1566–1581.
- Terui, N., Ban, M., & Allenby, G. M. (2011). The effect of media advertising on brand consideration and choice. *Marketing Science*, 30(1), 74–91.
- Tierney, L. (1994). Markov chains for exploring posterior distributions (with Discussion). *The Annals of Statistics*, 22(4), 1701–1762.
- Tierney, L., & Kadane, J. (1986). Accurate approximations for posterior moments and marginal densities. *Journal of the American Statistical Association*, 81(393), 82–86.
- Tierney, L., Kass, R., & Kadane, J. (1989). Fully exponential Laplace approximations to expectations and variances of non-positive functions. *Journal of the American Statistical Association*, 84(407), 710–716.
- Toubia, O., Iyengar, G., Bunnell, R., & Lemaire, A. (2019). Extracting features of entertainment products: A guided latent Dirichlet allocation approach informed by the psychology of media consumption. *Journal of Marketing Research*, 56(1), 18–36.
- Train, K. E. (2009). *Discrete choice methods with simulation*. Cambridge University Press.
- Van Gael, J., Saatci, Y., Teh, Y. W., & Ghahramani, Z. (2008). Beam sampling for the infinite hidden Markov model. In *Proceedings of the 25th international conference on machine learning* (pp. 1088–1095). ACM.
- Vankov, E. R., Guindani, M., & Ensor, K. B. (2019). Filtering and estimation for a class of stochastic volatility models with intractable likelihoods. *Bayesian Analysis*, 14(1), 29–52.
- Virbickaitė, A., Ausín, M. C., & Galeano, P. (2020). Copula stochastic volatility in oil returns: Approximate Bayesian computation with volatility prediction. *Energy Economics*, 92, Article 104961.
- Walker, S. G. (2007). Sampling the Dirichlet mixture model with slices. *Communications in Statistics—Simulation and Computation*, 36(1), 45–54.
- Wand, M. P. (2017). Fast approximate inference for arbitrarily large semiparametric regression models via message passing. *Journal of the American Statistical Association*, 112(517), 137–168.
- Wang, X., Hyndman, R. J., Li, F., & Kang, Y. (2022). Forecast combinations: An over 50-year review. *International Journal of Forecasting*, <https://doi.org/10.1016/j.ijforecast.2022.11.005>.
- Wang, S., Sun, X., & Lall, U. (2017). A hierarchical Bayesian regression model for predicting summer residential electricity demand across the USA. *Energy*, 140(1), 601–611.
- Weron, R. (2014). Electricity price forecasting: A review of the state-of-the-art with a look into the future. *International Journal of Forecasting*, 30(4), 1030–1081.
- West, M., & Harrison, J. (2006). *Bayesian forecasting and dynamic models*. Springer Science & Business Media.
- Wood, S. (2019). Simplified integrated nested Laplace approximation. *Biometrika*, 107(1), 223–230.
- Yang, Q. (2019). Stock returns and real growth: A Bayesian nonparametric approach. *Journal of Empirical Finance*, 53, 53–69.
- Yang, Y., Li, W., Gulliver, T. A., & Li, S. (2019). Bayesian deep learning-based probabilistic load forecasting in smart grids. *IEEE Transactions on Industrial Informatics*, 16(7), 4703–4713.
- Yu, C. L., Li, H., & Wells, M. T. (2011). MCMC estimation of Levy jump models using stock and option prices. *Mathematical Finance*, 21(3), 383–422.
- Zaharieva, M. D., Trede, M., & Wilfling, B. (2020). Bayesian semi-parametric multivariate stochastic volatility with application. *Econometric Reviews*, 39(9), 947–970.
- Zamenjani, A. S. (2021). Do financial variables help predict the conditional distribution of the market portfolio? *Journal of Empirical Finance*, 62, 327–345.
- Zellner, A. (1971). *An Introduction to Bayesian Econometrics*. New York: John Wiley.
- Zhang, C., Büttepage, J., Kjellström, H., & Mandt, S. (2018). Advances in variational inference. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 41(8), 2008–2026.
- Zhang, W., Smith, M. S., Maneesoonthorn, W., & Loaiza-Maya, R. (2023). Natural gradient hybrid variational inference with application to deep mixed models. arXiv preprint arXiv:2302.13536.
- Zhou, X., Nakajima, J., & West, M. (2014). Bayesian forecasting and portfolio decisions using dynamic dependent sparse factor models. *International Journal of Forecasting*, 30(4), 963–980.